

---

# Database Systems

## Lecture #3

---



# Agenda

- Last time:
  - A little on design
  - (nearly) finished E/R models
- This time:
  - Finish E/R
  - Constraints (some review)
  - Relational model
  - Converting E/R to relations
- Next time: Functional dependencies (Chapter19)




# Quick topic: Weak entity sets

- Def: some or all key attributes belong to another ES
- Plays role in a connecting relationship
- The key consists of:
  - Possibly its own attributes and
  - All key attributes of entity sets from supporting relationships



# Conditions for supporting relationships




- Supporting relationship  $R: E \rightarrow F$ 
  - $R$  is many-one (or 1-1)  $E \rightarrow F$
  - With referential integrity (rounded arrow)
  - $R$  is binary
  - $E$  receives key attributes of  $F$
- $F$  itself may be weak
  - Another entity set  $G$ , and so on recursively...



# Conditions for weak entity sets

- For several supporting relships from E to F
  - Keys of each F role appear as foreign key of E



- Other, non-supporting many-one relationships are not affected



# Weak entity set e.g.

- Example: Hierarchy – species & genus
- Idea: species name unique *per genus only*



- 物种(Species)名称一样，只能根据基因(Genus)不同区分




# Next topic: Constraints (约束)

- Review: programmer-defined rules stating what should always be true about consistent databases
- Restrictions on data (egs?):
  - ❑ Keys
  - ❑ Single value constraints
  - ❑ Referential Integrity (参照完整性)
  - ❑ Domain constraints
  - ❑ General constraints
- Can't infer (推断) constraints from data
  - ❑ may hold “accidentally”
  - ❑ but they are **a part of the schema**



# E/R keys

- Uniquely identify entity in ES
- Attribute or *set of* attributes
  - Two entities cannot agree on all key attributes
  - These attributes determine all others
- Every ES should have a key
  - possibly including all attributes
- Primary key attributes underlined
- More than one possible key:
  - *Candidate keys, primary key*
- Practical tip: create art key attribute
  - E.g. ID, course-id, employee-id, etc.
  - ID shorter than (name,address)





# Single-valued constraints

- “at most one” value
  - Already saw sharp arrows for relationships




- Attributes have this automatically
  - could be null or one value
  - Can think of key atts as (non-null) single-valued

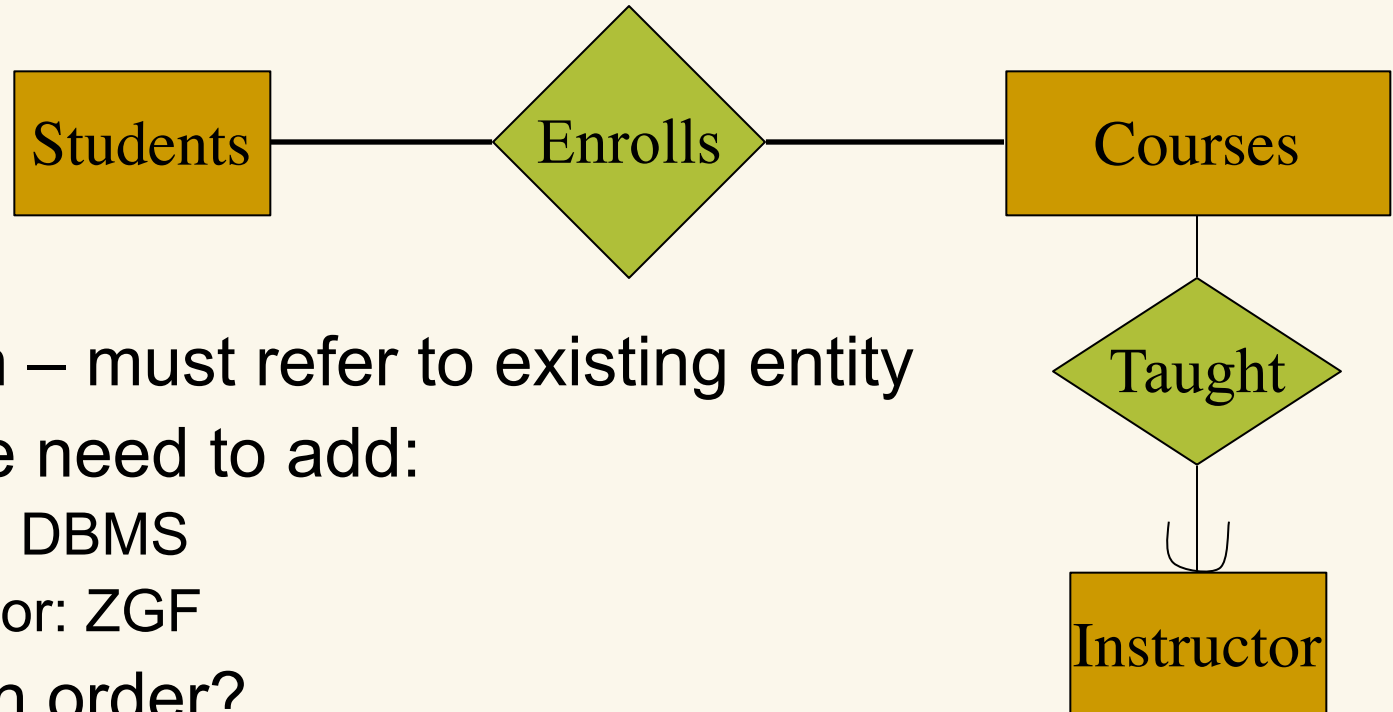


# Referential integrity

- “Exactly one value”
  - ❑ NOT NULL & *foreign keys* in SQL
- Relationships
  - ❑ Non-null value refers to entity that exists
  - ❑ Refer to entity with foreign key
  - ❑ HTML analogy: no broken links
  - ❑ Programming analogy: no dangling pointers
  - ❑ Multiple ways of handling violations...



# Referential integrity – E/R e.g.




- Insertion – must refer to existing entity
- Suppose need to add:
  - course: DBMS
  - instructor: ZGF
- Q: Which order?
- Q: What if relship were exactly-exactly, say, M(Hs,Ws)?
  - i.e., referential integrity in both directions?
- A: Put both inserts in one xact – later




# Other kinds of constraints

- Domain constraints
  - ❑ E.g. date: must be after 1980
  - ❑ Enumerated type: grades A through F, no E
  - ❑ No special E/R notation just write near line
- General constraints:
  - ❑ A class may have no more than 100 students; a student may not have more than 6 courses:





# Next topic: the Relational Data Model

- Invented by Ted Codd
  - Researcher at IBM
  - We'll see his name again...
- Related work at Berkeley
- Introduced in a paper a paper published in June, 1970



# DB development path



# Relations as tables

Attribute  
names

Product table/relation

Name	Price	Category	Manufacturer
gizmo	\$19.99	gadgets	GizmoWorks
Power gizmo	\$29.99	gadgets	GizmoWorks
SingleTouch	\$149.99	photography	Canon
MultiTouch	\$203.99	household	Hitachi

tuples/rows/records/entities



# Relational terminology

- Relation is composed of *tuples*
- Tuple = sequence of *attribute* values
  - Attribute has atomic types
- Relation *schema*:  
relation name + attribute names + attribute types
- Database schema: set of relation schemas





# Relations as sets

- Recall: math relation is a subset of the cross-product of the attribute value sets
  - $R \text{ subset-of } S \times T$
  - Product subset-of Name x Price x Cat x Mft
- One member of Product relation:
  - (gizmo, \$19.99, gadgets, GizmoWorks) in Product
  - $\text{Product}(\text{gizmo}, \$19.99, \text{gadgets}, \text{GizmoWorks})$
- Usual updates: add/delete/change a tuple in this set
- Updates to the *schema* are rare, painful (why?)



# From E/R models to relations

- Recall justification:
  - ❑ design is easier in E/R
  - ❑ but implementation is easier/faster in R
- Analogy to program compilation:
  - ❑ design is easier in C/Java/whatever
  - ❑ implemen. is easier/faster in machine/byte code
- Strategy:
  1. apply semi-mechanical conversion rules
  2. improve by combining some relations
  3. improve by normalization
    - involves finding functional dependencies




# E/R conversion rules

- Relationship → relation
  - attributes: keys of entity-sets/roles
  - key: depends on multiplicity
- Entity set → ... relation
  - attributes: attributes of entity set
  - key: key of ES
- NB: mapping of types is not one-one
  - We'll see: mapping one tokens is also not one-one
- Special treatment:
  - Weak entity sets
  - Isa relations & subclasses




# Entity Sets

- Entity set Students



Rel: Students

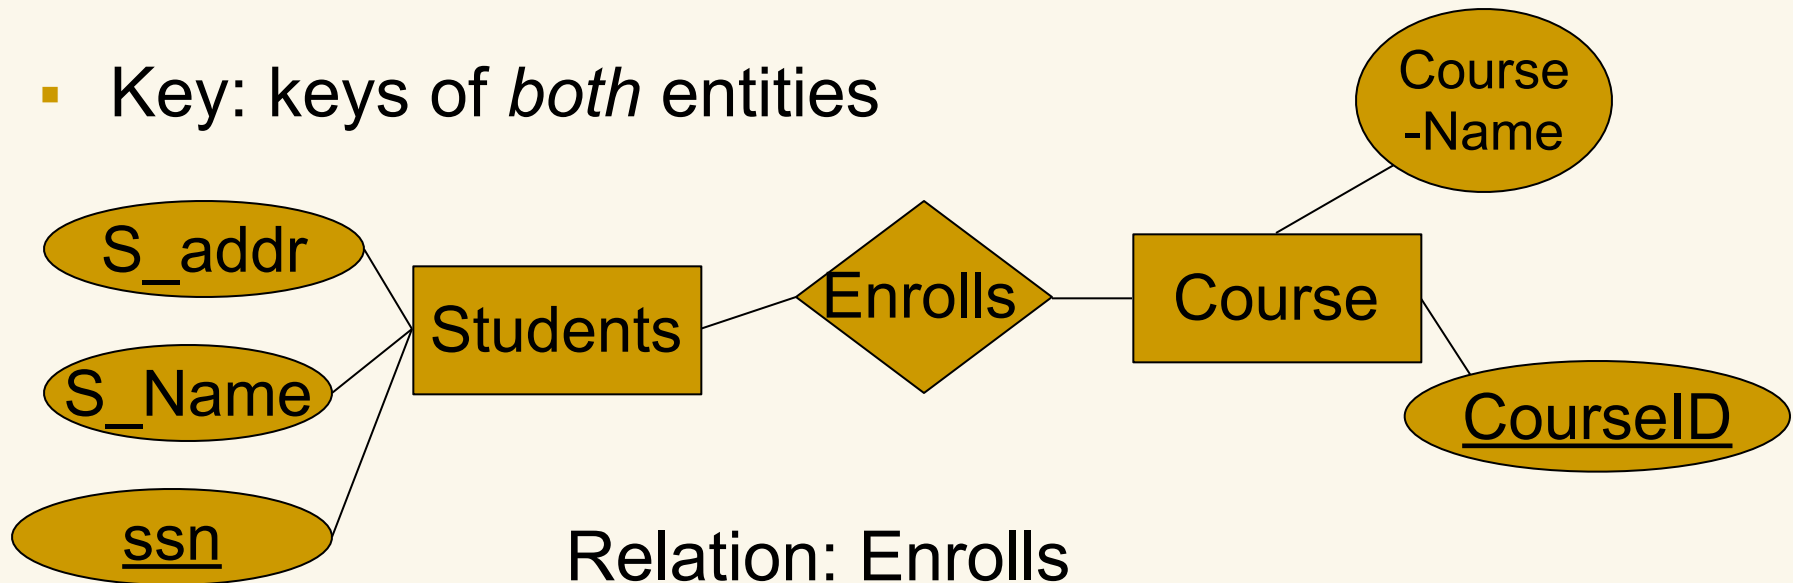


<u>SSN</u>	Name	Address
111-222-3333	Howard	Park Avenue
444-555-6666	John	South Carolina




# Binary many-to-many relationships

- Key: keys of *both* entities




Relation: Enrolls



<u>ssn</u>	<u>CourseID</u>
111-222-3333	SE-304
111-222-3333	C20.0056
444-555-6666	C30.0046



# Many-to-one relationships



Movies

<u>MovieID</u>	Title	Year
M101	Mr. Ripley.	1999
M202	Sylia	2003

Studios

<u>StudioID</u>	Name	Address
S35	Miramax	NYC
S73	Disney	Orlando

- Key: keys of *many* entity

Owns

<u>MovieID</u>	StudioID	CopyrightNo
M101	S73	CN11111
M202	S35	CN22222



# Many-to-one: a better design


Movies

<u>MovieID</u>	Title	Year
M101	Mr. Ripley.	1999
M202	Sylia	2003

Owns

<u>MovieID</u>	StudioID	CopyrightNo
M101	S73	CN11111
M202	S35	CN22222

Movies'



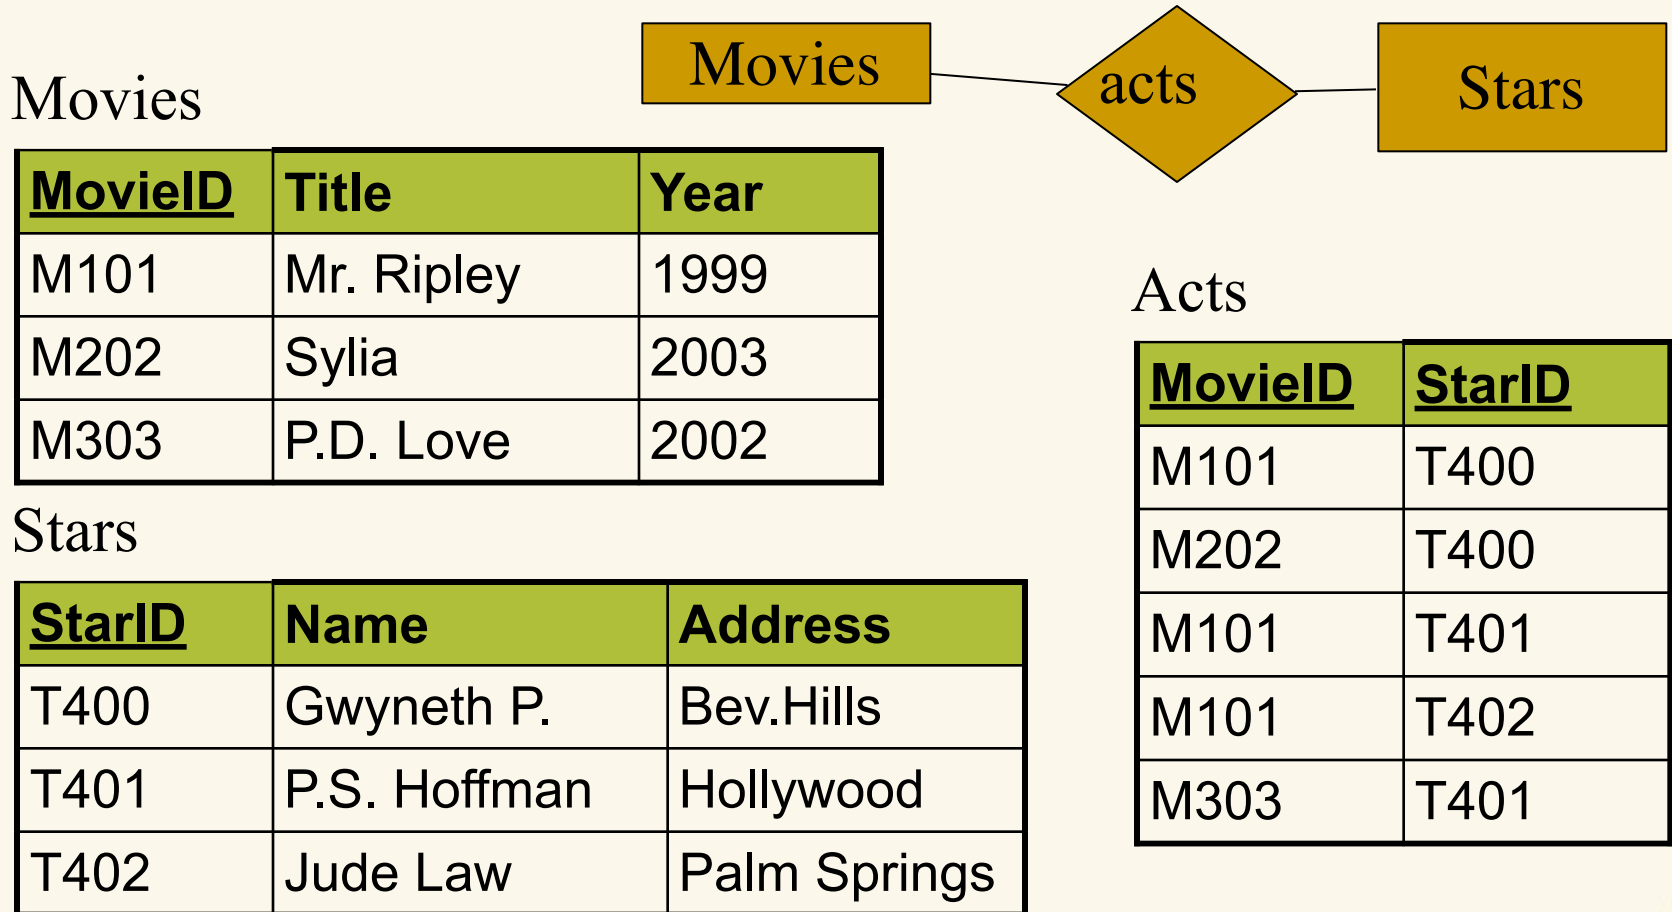
<u>MovieID</u>	Title	Year	StudioID	CopyrightNo
M101	Talent Mr. Ripley	1999	S73	CN11111
M202	Sylia	2003	S35	CN22222

- Q: What if a movie's Owns row were missing?



# Many-to-many relationships again

- NB: Won't work for many-many relationships





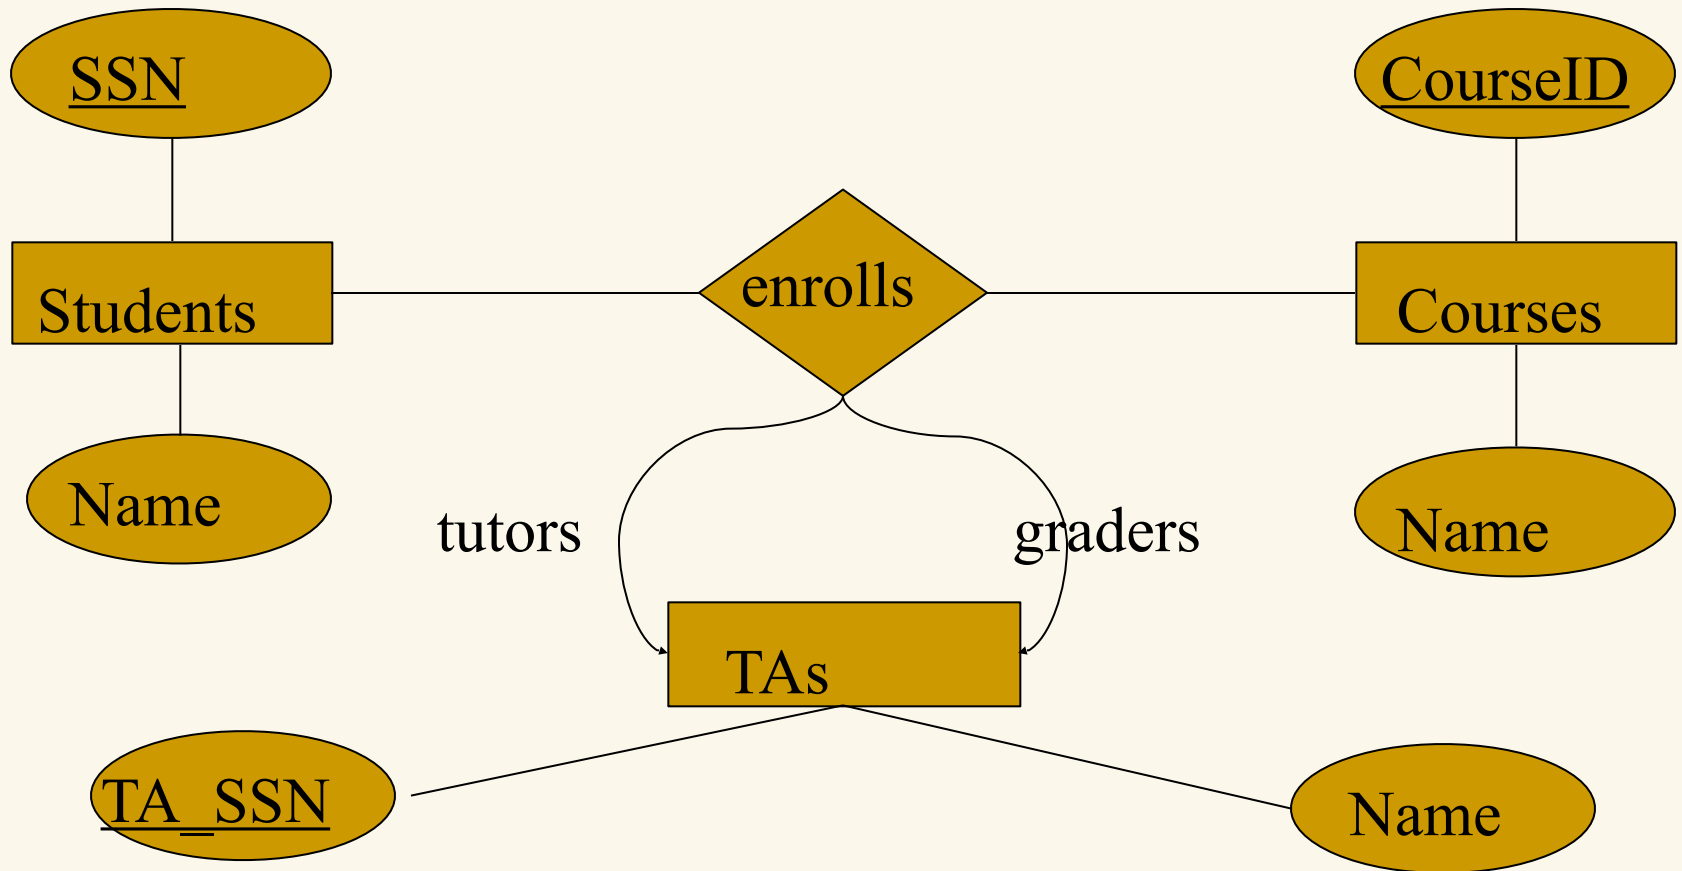
# Many-to-many relationships again

And here's why:

<b>MovieID</b>	<b>Title</b>	<b>Year</b>	<b>StarID</b>
M101	Talented Mr. Ripley	1999	T400
M101	Talented Mr. Ripley	1999	T401
M101	Talented Mr. Ripley	1999	T402
M202	Sylvia	2003	T400
M303	Punch Drunk Love	2003	T401



# Multiway relationships & roles



- Different roles treated as different entity sets
- Key: keys of the *many* entities



# Multiway relationships & roles

Students

<u>SSN</u>	Name
111-11-1111	George
222-22-2222	Dick

TAs

<u>TA_SSN</u>	Name
333-33-3333	Wesley
444-44-4444	Howard
555-55-5555	John

Courses


<u>CourseID</u>	Name
SE-304	Databases
C20.0056	Software

Enrolls(S\_SSN, Course\_ID, Tutor\_SSN, Grader\_SSN)

<u>S_SSN</u>	<u>CourseID</u>	Tutor_SSN	Grader_SSN
111-11-1111	SE-304	333-33-3333	444-44-4444
222-22-2222	SE-304	444-44-4444	555-55-5555



# Converting weak ESs – differences



- Atts of Jersey Rel. are:
  - attributes of Jersey
  - key attributes of supporting ESs, Club

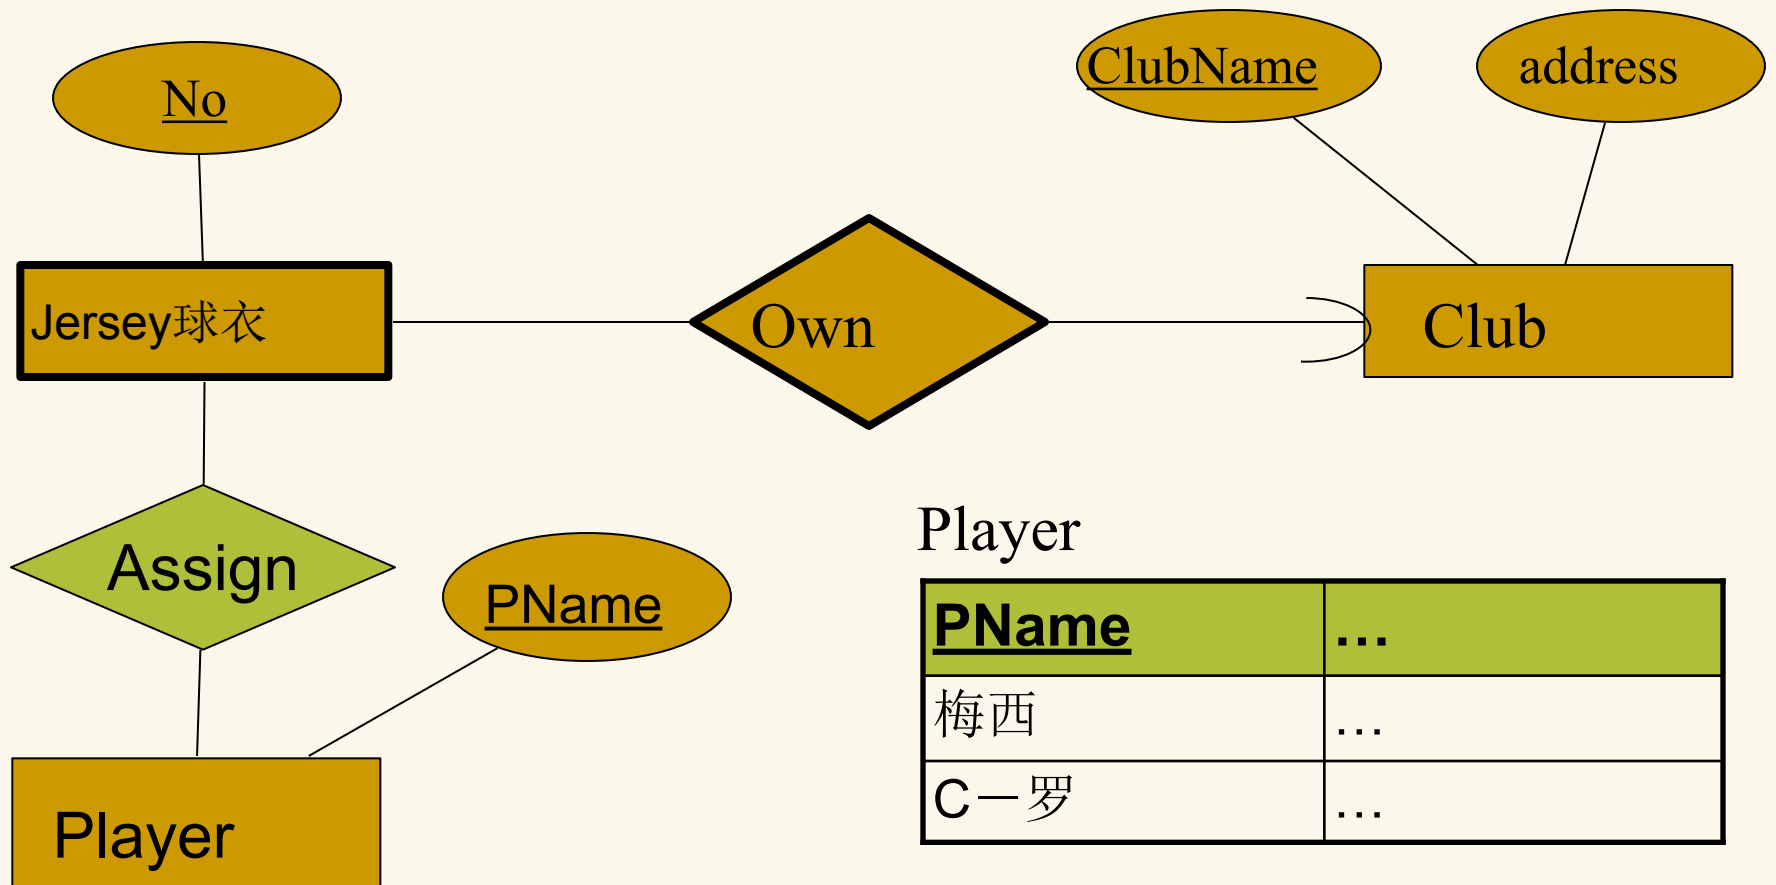
Jersey

<u>ClubName</u>	<u>JerseyNo</u>
巴萨	10
巴萨	9
皇马	10

- Supporting relships are omitted (*why?*)

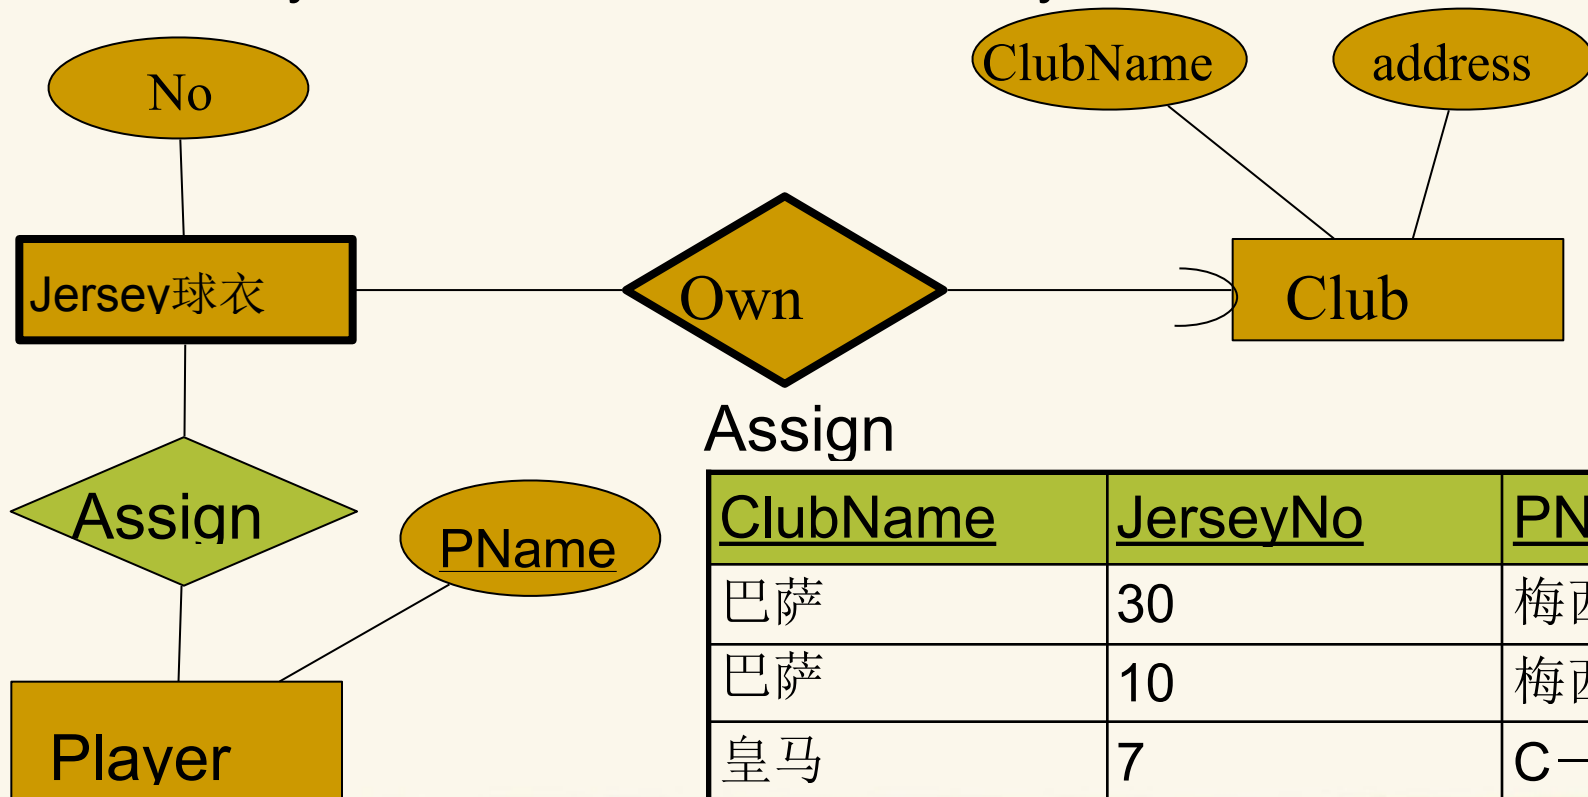


# Weak entity sets - relationships



# Weak entity sets - relationships

- Non-supporting relationships for weak ESs *are* converted
  - keys include entire weak ES key



# Next week

- For next week:
  - Review/skim Ch.3 section 5 (from today)
  - Read Ch.19 sections 1-3

