

Relational Algebra

关系代数

courtesy of Joe Hellerstein for some slides

Guifeng Zheng

SUN YAT-SEN UNIVERSITY

Recap: You are here

- First part of course is done: conceptual foundations
- You now know:
 - ❑ E/R Model
 - ❑ Relational Model
 - ❑ Relational Algebra (a little, project / join)
- You now know how to:
 - ❑ Capture part of world as an E/R model
 - ❑ Convert E/R models to relational models
 - ❑ Convert relational models to good (normal) forms
- Next:
 - ❑ Create, update, query tables with R.A/SQL
 - ❑ Write SQL/DB-connected applications

3-minute Normalization Review

Q: What's required for BCNF?

Q: How do we fix a non-BCNF relation?

Q: If $A \rightarrow B$ violates BCNF, what do we do?

Q: Can BCNF decomposition ever be lossy?

Q: How do we combine two relations?

Q: Can BCNF decomp. lose FDs?

Q: Why would you ever use 3NF?

Relational Query Languages

- Query languages: manipulation and retrieval of data
- Relational model supports simple, powerful QLs:
 - Strong formal foundation based on logic.
 - Allows for much optimization.
- Query Languages **!=** programming languages!
 - QLs not expected to be “Turing complete”.
 - QLs not intended to be used for complex calculations.
 - QLs support easy, efficient access to large data sets.

(Actually, I no longer believe this. But it's the standard viewpoint)

Formal Relational Query Languages

Relational Algebra关系代数:

More **operational**, very useful for representing execution plans.

Relational Calculus关系演算:

Describe what you want, rather than how to compute it. (**Non-procedural**, declarative.)

👉 *Understanding Algebra & Calculus is key to understanding SQL, query processing!*

What is relational algebra?

- An algebra for relations
- “High-school” algebra: an algebra for *numbers*
- *Algebra* = formalism for constructing expressions
 - Operations
 - Operands: Variables, Constants, expressions
- Expressions:
 - Vars & constants
 - Operators applied to expressions
 - They evaluate to values

Algebra	Vars/consts	Operators	Eval to
High-school	Numbers	+ * - / etc.	Numbers
Relational	Relations (=sets of tupes)	union, intersection, join, etc.	Relations

Why do we care about relational algebra?

- ❑ The exprs are *the form that questions about the data take* （有关数据的问题采用的形式！）
 - ❑ The relations these exprs cash out to are *the answers to our questions* （其表示的关系正是我们的问题的答案）
- ❑ RA ~ more succinct rep.(简洁表示) of many SQL queries
- ❑ DBMS parse SQL into something like RA.
- First proofs of concept for RDBMS/RA:
 - ❑ System R at IBM
 - ❑ Ingress at Berkeley
- “Modern” implementation of RA: SQL
 - ❑ Both state of the art, mid-70s

Preliminaries 预备知识

- A query is applied to *relation instances*
- The result of a query is also a relation instance.
 - *Schemas* of input relations for a query are *fixed*
 - *Schema* for the *result* of a query is also *fixed*.
 - determined by the query language constructs
- Positional vs. named-field notation:
 - Positional notation easier for formal definitions
 - Named-field notation more readable.
 - Both used in SQL
 - Though positional notation is discouraged

Relational Algebra: 5 Basic Operations

- **Selection** (σ)
 - ❑ Selects a subset of **rows** (horizontal)
- **Projection** (π)
 - ❑ Retains only desired **columns** (vertical)
- **Cross-product** (\times)
 - ❑ Allows us to combine two relations.
- **Set-difference** ($-$)
 - ❑ Tuples in r1, but **not** in r2.
- **Union** (\cup)
 - ❑ Tuples in r1 **or** in r2.
- ❑ Since each operation returns a relation, **operations can be composed!** (Algebra is “closed”.)

Example Instances

Boats

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

Projection (π)

- **Example:** $\pi_{sname, rating}(S2)$
- **Retains only attributes that are in the “*projection list*”.**
- ***Schema* of result:**
 - the fields in the projection list
 - with the same names that they had in the input relation.
- **Projection operator has to *eliminate duplicates***
 - Note: real systems typically don't do duplicate elimination
 - Unless the user explicitly asks for it.
 - (Why not?)

Projection (π)

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S2

sname	rating
yuppy	9
lubber	8
guppy	5
rusty	10

$\pi_{sname, rating}(S2)$

age
35.0
55.5

$\pi_{age}(S2)$

Selection (σ)

- **Selects rows that satisfy *selection condition*.**
- **Result is a relation.**
Schema of result is same as that of the input relation.
- **Do we need to do duplicate elimination?**

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

$\sigma_{rating > 8}(S2)$

sname	rating
yuppy	9
rusty	10

$\pi_{sname, rating}(\sigma_{rating > 8}(S2))$

Union, Set-Difference

- **Both of these operations take two input relations, which must be union-compatible:**
 - Same number of fields.
 - 'Corresponding' fields have the same type.
- **For which, if any, is duplicate elimination required?**

Union

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S2

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0
44	guppy	5	35.0
28	yuppy	9	35.0

$S1 \cup S2$

Set Difference

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S2

sid	sname	rating	age
22	dustin	7	45.0

$S1 - S2$

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
44	guppy	5	35.0

$S2 - S1$

Cross-Product

- **$S1 \times R1$:**
 - Each row of **S1** paired with each row of **R1**.
- **Q: How many rows in the result?**
- ***Result schema* has one field per field of S1 and R1,**
 - Field names 'inherited' if possible.

Cross Product Example

S1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

R1

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

S1 x R1 =

Naming conflict: S1
and R1 have a field
with the same name.

(Can use the
renaming operator)

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	22	101	10/10/96
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	22	101	10/10/96
31	lubber	8	55.5	58	103	11/12/96
58	rusty	10	35.0	22	101	10/10/96
58	rusty	10	35.0	58	103	11/12/96

Rename op

- Changes the *schema*, not the *instance*
- Notation: $\rho_{B1, \dots, Bn}(R)$
- ρ is spelled “rho”, pronounced “row”

$$\rho_{(C(1 \rightarrow \text{sid1}, 5 \rightarrow \text{sid2})}(R1XS1)$$

- Example:
 - ❑ Employee(ssn, name)
 - ❑ $\rho_{E2(\text{social}, \text{name})}(\text{Employee})$
 - ❑ Or just: $\rho_E(\text{Employee})$

Compound Operator: Intersection

- **On top of 5 basic operators, several additional “Compound Operators”**
 - These add no computational power to the language
 - Useful shorthand
 - Can be expressed solely with the basic operators.
- **Intersection takes two input relations, which must be union-compatible.**
- **Q: How to express it using basic operators?**

$$\mathbf{R \cap S = R - (R - S)}$$

Intersection

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

<u>sid</u>	sname	rating	age
28	yuppy	9	35.0
31	lubber	8	55.5
44	guppy	5	35.0
58	rusty	10	35.0

S2

<u>sid</u>	sname	rating	age
31	lubber	8	55.5
58	rusty	10	35.0

$S1 \cap S2$

Compound Operator: Join

- Involve cross product, selection, and (sometimes) projection.
- Most common type of join: “natural join”
 - $R \bowtie S$ conceptually is:
 - Compute $R \times S$
 - Select rows where attributes appearing in both relations have equal values
 - Project all unique attributes and one copy of each of the common ones.
- Note: Usually done much more efficiently than this.

Natural Join Example

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

R1

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

S1

S1 ⋈ R1 =

sid	sname	rating	age	bid	day
22	dustin	7	45.0	101	10/10/96
58	rusty	10	35.0	103	11/12/96

Natural Join

- R

A	B
X	Y
X	Z
Y	Z
Z	V

- S

B	C
Z	U
V	W
Z	V

- $R \bowtie S = ?$

- Unpaired tuples called *dangling*

Natural Join

- Given the schemas $R(A, B, C, D)$, $S(A, C, E)$, what is the schema of $R \bowtie S$?
- Given $R(A, B, C)$, $S(D, E)$, what is $R \bowtie S$?
- Given $R(A, B)$, $S(A, B)$, what is $R \bowtie S$?

Other Types of Joins

- Condition Join (or “theta-join”):

$$R \bowtie_c S = \sigma_c (R \times S)$$

(sid)	sname	rating	age	(sid)	bid	day
22	dustin	7	45.0	58	103	11/12/96
31	lubber	8	55.5	58	103	11/12/96

$$S1 \bowtie_{S1.sid < R1.sid} R1$$

- **Result schema** same as that of cross-product.
- **May have fewer tuples than cross-product.**
- Equi-Join: Special case: condition c contains only conjunction of **equalities**.

Division Operation

- Notation: $r \div s$
- Suited to queries that include the phrase “for all.”
- Let r and s be relations over schemas R and S respectively, where

$$R = (A_1, \dots, A_m, B_1, \dots, B_n)$$

$$S = (B_1, \dots, B_n)$$

The result of $r \div s$ is a relation over the schema $(R - S) = (A_1, \dots, A_m)$

$$r \div s = \{ t \mid (t \in \pi_{R-S}(r)) \wedge (\forall u \in s, tu \in r) \}$$

Division Operation - example

$$r \div s = \{ t \mid (t \in \pi_{R-S}(r)) \wedge (\forall u \in s, tu \in r) \}$$

r

A	B
α	1
α	2
α	3
β	1
γ	1
δ	1
δ	3
δ	4
δ	6
ε	1
ε	2

s

B
1
2

The result consists of attribute A only but not all of the 5 values. How to find out?
 $u = 1, 2$ Check if: $\forall u \in s (tu \in r)$

$$t \in \pi_{R-S}(r)$$

A
α
β
γ
δ
ε

Is $\langle \alpha, 1 \rangle$ and $\langle \alpha, 2 \rangle$ tuples in r?

Is $\langle \beta, 1 \rangle$ and $\langle \beta, 2 \rangle$ tuples in r?

check γ and δ ...

Is $\langle \varepsilon, 1 \rangle$ and $\langle \varepsilon, 2 \rangle$ tuples in r?

$r \div s$

A
α
ε

Another Division Example

Relations r, s:

A	B	C	D	E
α	A	α	A	1
α	A	γ	A	1
α	A	γ	B	1
β	A	γ	A	1
β	A	γ	B	3
γ	A	γ	A	1
γ	A	γ	B	1
γ	A	β	B	1

D	E
A	1
B	1

$r \div s$

A	B	C
α	A	γ
γ	A	γ

Properties of Division Operation

- Let $q = r \div s$

Then q is **the largest** relation satisfying: $q \times s \subseteq r$

Relation r, s :

r

A	B
α	1
α	2
α	3
β	1
γ	1
δ	1
δ	3
δ	4
δ	6
ε	1
ε	2

s

B
1
2

q

A
α
ε

Examples

Reserves

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
58	103	11/12/96

Sailors

<u>sid</u>	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.5
58	rusty	10	35.0

Boats

<u>bid</u>	bname	color
101	Interlake	Blue
102	Interlake	Red
103	Clipper	Green
104	Marine	Red

Find names of sailors who've reserved boat #103

- Solution 1:

$$\pi_{name} ((\sigma_{bid=103} \text{Reserves}) \bowtie \text{Sailors})$$

- Solution 2:

$$\pi_{name} (\sigma_{bid=103} (\text{Reserves} \bowtie \text{Sailors}))$$

Find names of sailors who've reserved a red boat

- Information about boat color only available in Boats; so need an extra join:

$$\pi_{sname}((\sigma_{color='red'} Boats) \bowtie Reserves \bowtie Sailors)$$

❖ A more efficient solution:

$$\pi_{sname}(\pi_{sid}((\pi_{bid} \sigma_{color='red'} Boats) \bowtie Res) \bowtie Sailors)$$

👉 A query optimizer can find this given the first solution!

Find sailors who've reserved a red or a green boat

- Can identify all red or green boats, then find sailors who've reserved one of these boats:

$$\rho \text{ (Tempboats, } (\sigma_{color='red' \vee color='green'} \text{Boats}))$$
$$\pi_{sname}(\text{Tempboats} \bowtie \text{Reserves} \bowtie \text{Sailors})$$

Find sailors who've reserved a red and a green boat

- Cut-and-paste previous slide?

$\rho (Tempboats, (\sigma_{color='red' \wedge color='green'} Boats))$

$\pi_{sname}(Tempboats \bowtie Reserves \bowtie Sailors)$



Find sailors who've reserved a red and a green boat

- Previous approach won't work! Must identify sailors who've reserved red boats, sailors who've reserved green boats, then find the intersection (note that *sid* is a key for *Sailors*):

$$\rho \text{ (Tempred, } \pi_{sid}((\sigma_{color='red'} Boats) \bowtie Reserves))$$
$$\rho \text{ (Tempgreen, } \pi_{sid}((\sigma_{color='green'} Boats) \bowtie Reserves))$$
$$\pi_{sname}((Tempred \cap Tempgreen) \bowtie Sailors)$$

Summary

- Relational Algebra: a small set of operators mapping relations to relations
 - Operational, in the sense that you specify the explicit order of operations
 - A *closed* set of operators! Can mix and match.
- Basic ops include: σ , π , \times , \cup , $-$,
- Important compound ops: \cap , \bowtie