

# SQL: The Query Language

Guifeng Zheng  
SUN YAT-SEN UNIVERSITY

The important thing is not to  
stop questioning.

Albert Einstein

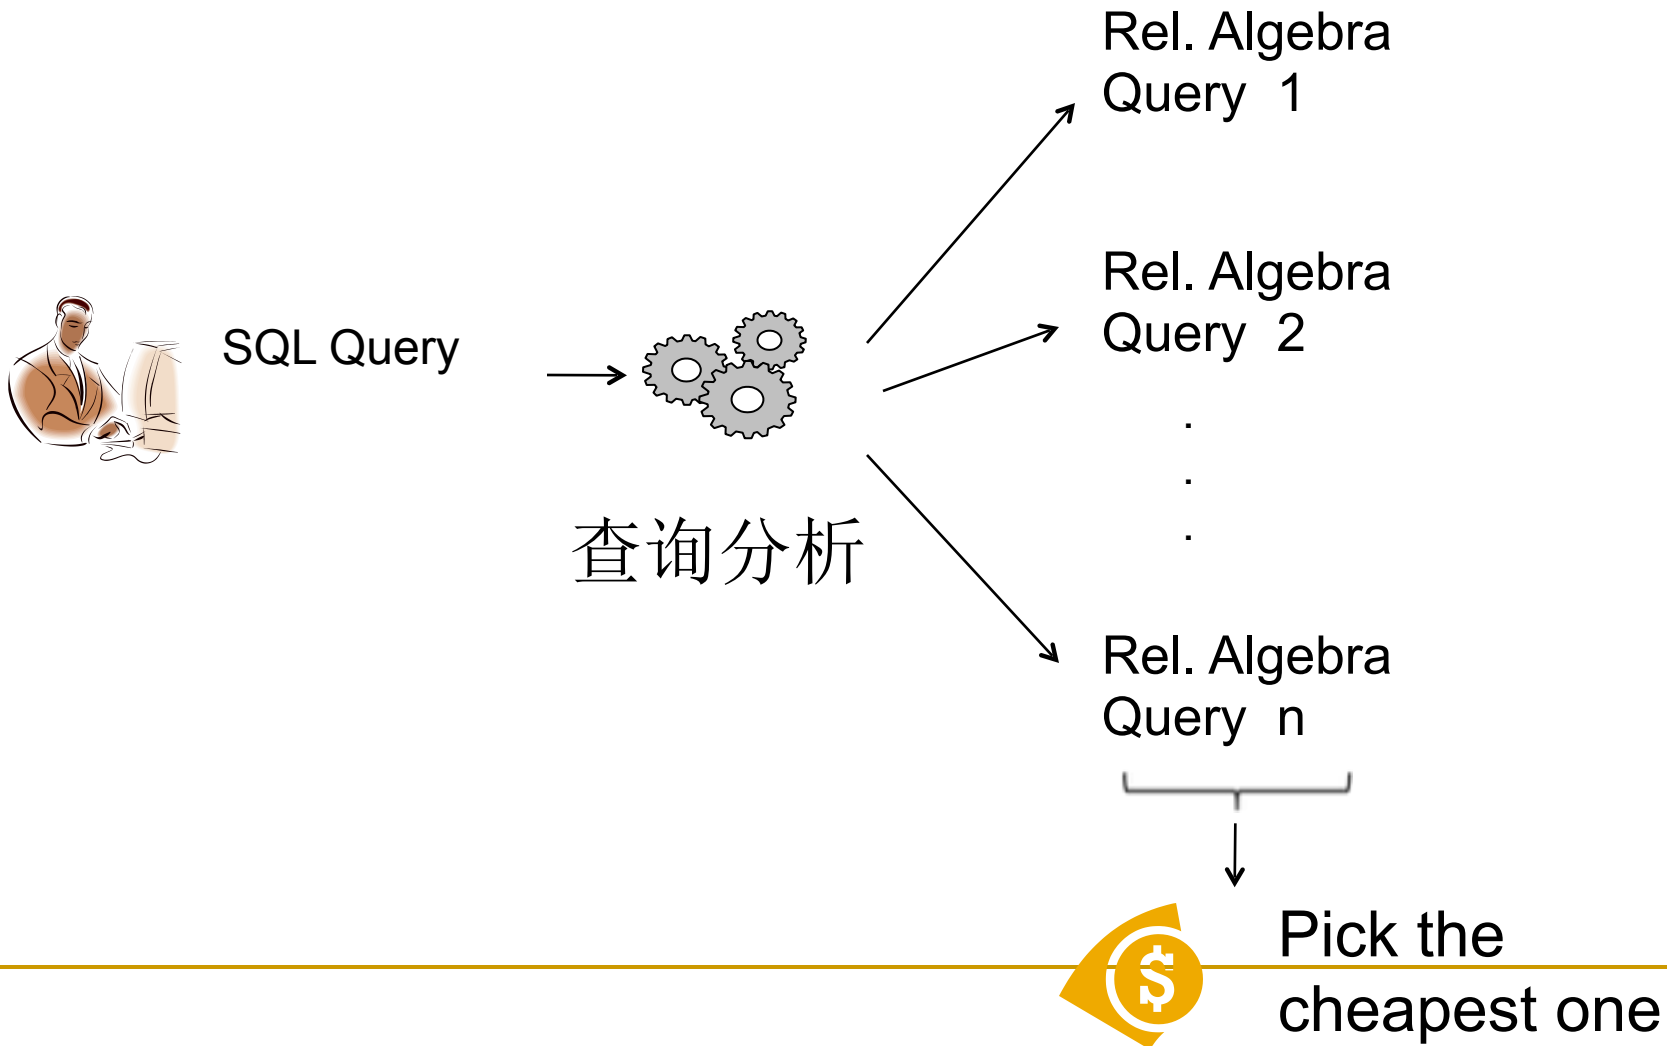
# Review

- Relational Algebra (Operational Semantics操作语义)
  - Given a query, how to mix and match the relational algebra operators to answer it
  - Used for query optimization用于查询优化
- Relational Calculus (Declarative Semantics说明性语义)
  - Given a query, what do I want my answer set to include?
- Algebra and safe calculus are simple and powerful models for query languages for relational model
  - Have same expressive power有相同的表达力
- SQL can express every query that is expressible in relational algebra/calculus. (and more)

# RA $\rightarrow$ SQL

- SQL SELECT  $\rightarrow$  RA Projection  $\Pi$
- SQL WHERE  $\rightarrow$  RA Selection  $\sigma$
- SQL FROM  $\rightarrow$  RA Join/cross
  - Comma-separated list...
- SQL renaming  $\rightarrow$  RA rho  $\rho$
- More ops later
- *Keep RA in the back of your mind...*

# Relational Query Languages



# Relational Query Languages

- Two sublanguages:
  - DDL – Data Definition 定义 Language
    - Define and modify schema (at all 3 levels)
  - DML – Data Manipulation 操作 Language
    - Queries can be written intuitively.
- DBMS is responsible for efficient evaluation.
  - The key: precise semantics for relational queries.
  - Optimizer can re-order operations
    - Won't affect query answer.
  - Choices driven by “cost model” 成本模型

---

# The SQL Query Language

- The most widely used relational query language.
  - Standardized  
(although most systems add their own “special sauce” -- including PostgreSQL)
  - We will study SQL92 -- **a basic subset**
-

# Example Database

## Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

## Boats

bid	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

## Reserves

sid	bid	day
1	102	9/12
2	102	9/13

# The SQL DDL

```
CREATE TABLE Sailors (  
  sid INTEGER,  
  sname CHAR(20),  
  rating INTEGER,  
  age REAL,  
  PRIMARY KEY sid);
```

```
CREATE TABLE Boats (  
  bid INTEGER,  
  bname CHAR(20),  
  color CHAR(10),  
  PRIMARY KEY bid);
```


```
CREATE TABLE Reserves (  
  sid INTEGER,  
  bid INTEGER,  
  day DATE,  
  PRIMARY KEY (sid, bid, day),  
  FOREIGN KEY sid REFERENCES Sailors,  
  FOREIGN KEY bid REFERENCES Boats);
```

<u>sid</u>	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

<u>bid</u>	bname	color
101	Nina	red
102	Pinta	blue
103	Santa Maria	red

<u>sid</u>	<u>bid</u>	<u>day</u>
1	102	9/12
2	102	9/13

FOREIGN  
KEY 外键





# The SQL DML

## Sailors

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

- Find all 18-year-old sailors:

$\sigma_{\text{age}=18}(\text{Sailors})$

```
SELECT *  
FROM sailors S  
WHERE S.age=18
```

- To find just names and ratings, replace the first line:

$\pi_{\text{sname}, \text{rating}}(\sigma_{\text{age}=18}(\text{Sailors}))$ 

```
SELECT S.sname, S.rating
```

# Querying Multiple Relations

Cross Product

```
SELECT S.sname
FROM   Sailors S, Reserves R
WHERE  S.sid=R.sid AND R.bid=102
```

Natural Join

**Sailors**

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

**Reserves**

sid	bid	day
1	102	9/12
2	102	9/13

# Basic SQL Query

**DISTINCT**: optional. Answer should not contain duplicates.

SQL default: duplicates are not eliminated! (Result a “multiset”)

target-list : List of expressions over attributes of tables in *relation-list*

```
SELECT [DISTINCT] target-list
FROM      relation-list
WHERE qualification
```

The diagram shows the SQL query syntax in a box. Arrows point from the definitions to the corresponding parts of the query: from the DISTINCT definition to **[DISTINCT]**, from the target-list definition to *target-list*, from the relation-list definition to *relation-list*, and from the qualification definition to *qualification*.

qualification : Comparisons combined using AND, OR and NOT. Comparisons are Attr op const or Attr1 op Attr2, where op is one of =, <, >, ≠, etc.

relation-list : List of relation names, possibly with a *range-variable* after each name

# Query Semantics

SELECT	[DISTINCT]	<i>target-list</i>
FROM		<i>relation-list</i>
WHERE		<i>qualification</i>

1. FROM : compute cross product of tables.
2. WHERE : Check conditions, discard tuples that fail.
3. SELECT : Delete unwanted fields.
4. DISTINCT (*optional*) : eliminate duplicate rows.

**Note:** Probably the least efficient way to compute a query!

- ❑ *Query optimizer* will find more efficient ways to get the *same answer*.
-

# SQL Query Semantics

```
SELECT  a1, a2, ..., ak  
FROM    R1 AS x1, R2 AS x2, ..., Rn AS xn  
WHERE   Conditions
```

Parallel assignment – all tuples

```
Answer = {}  
for all assignments x1 in R1, ..., xn in Rn do  
    if Conditions then  
        Answer = Answer  $\cup$  { (a1,...,ak) }  
return Answer
```

Doesn't impose any order

# SQL Query Semantics

```
SELECT  a1, a2, ..., ak  
FROM    R1 AS x1, R2 AS x2, ..., Rn AS xn  
WHERE   Conditions
```

Nested loops:

```
Answer = {}  
for x1 in R1 do  
    for x2 in R2 do  
        ....  
        for xn in Rn do  
            if Conditions then  
                Answer = Answer  $\cup$  {(a1,...,ak)}  
return Answer
```

---

Find sailors who've reserved at least one boat

```
SELECT S.sid  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid
```

- Would DISTINCT make a difference here?
  - What is the effect of replacing *S.sid* by *S.sname* in the SELECT clause?
    - Would DISTINCT make a diff to this variant of the query?
-

# About Range Variables

- Needed when ambiguity could arise.
  - e.g., same table used multiple times in FROM (“self-join”)

```
SELECT  x.sname, x.age, y.sname, y.age
FROM    Sailors x, Sailors y
WHERE   x.age > y.age
```

**Sailors x**

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27

**Sailors y**

sid	sname	rating	age
1	Fred	7	22
2	Jim	2	39
3	Nancy	8	27



# Arithmetic Expressions

```
SELECT S.age, S.age-5 AS age1, 2*S.age AS age2
FROM   Sailors S
WHERE  S.sname = 'dustin'
```

```
SELECT S1.sname AS name1, S2.sname AS name2
FROM   Sailors S1, Sailors S2
WHERE  2*S1.rating = S2.rating - 1
```

# String Comparisons

```
SELECT S.sname  
FROM   Sailors S  
WHERE  S.sname LIKE 'B_%B'
```

`\_` stands for any **one** character and `%` stands for **0 or more** arbitrary characters.

Find sid's of sailors who've reserved a red or a green boat

```
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND
        (B.color='red' OR
         B.color='green')
```

... or:

```
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND
        B.color='red'
UNION
SELECT R.sid
FROM   Boats B, Reserves R
WHERE  R.bid=B.bid AND B.color='green'
```

Find sid's of sailors who've reserved a red and a green boat

```
SELECT R.sid  
FROM   Boats B,Reserves R  
WHERE  R.bid=B.bid AND  
       (B.color='red' AND B.color='green')
```

Find sid's of sailors who've reserved a red and a green boat

```
SELECT S.sid
FROM   Sailors S, Boats B, Reserves R
WHERE  S.sid=R.sid
        AND R.bid=B.bid
        AND B.color='red'
```

INTERSECT

```
SELECT S.sid
FROM   Sailors S, Boats B, Reserves R
WHERE  S.sid=R.sid
        AND R.bid=B.bid
        AND B.color='green'
```

Find sid's of sailors who've reserved a red and a green boat

- Could use a self-join:

```
SELECT R1.sid
FROM   Boats B1, Reserves R1,
       Boats B2, Reserves R2
WHERE  R1.sid=R2.sid
       AND R1.bid=B1.bid
       AND R2.bid=B2.bid
       AND (B1.color='red' AND B2.color='green')
```

Find sid's of sailors who have not reserved a boat

```
SELECT S.sid  
FROM   Sailors S
```

EXCEPT

```
SELECT S.sid  
FROM   Sailors S, Reserves R  
WHERE  S.sid=R.sid
```

# Nested Queries: IN

*Names of sailors who've reserved boat #103:*

```
SELECT S.sname
FROM   Sailors S
WHERE  S.sid IN
      (SELECT R.sid
       FROM   Reserves R
       WHERE  R.bid=103)
```



# Nested Queries: NOT IN


*Names of sailors who've not reserved boat #103:*

```
SELECT  S.sname
FROM    Sailors S
WHERE   S.sid NOT IN
        (SELECT  R.sid
         FROM    Reserves R
         WHERE   R.bid=103)
```

# Nested Queries with Correlation

*Names of sailors who've reserved boat #103:*

```
SELECT  S.sname
FROM    Sailors S
WHERE   EXISTS
        (SELECT *
         FROM Reserves R
         WHERE R.bid=103 AND S.sid=R.sid)
```



- Subquery must be recomputed for each Sailors tuple.
  - Think of subquery as a function call that runs a query

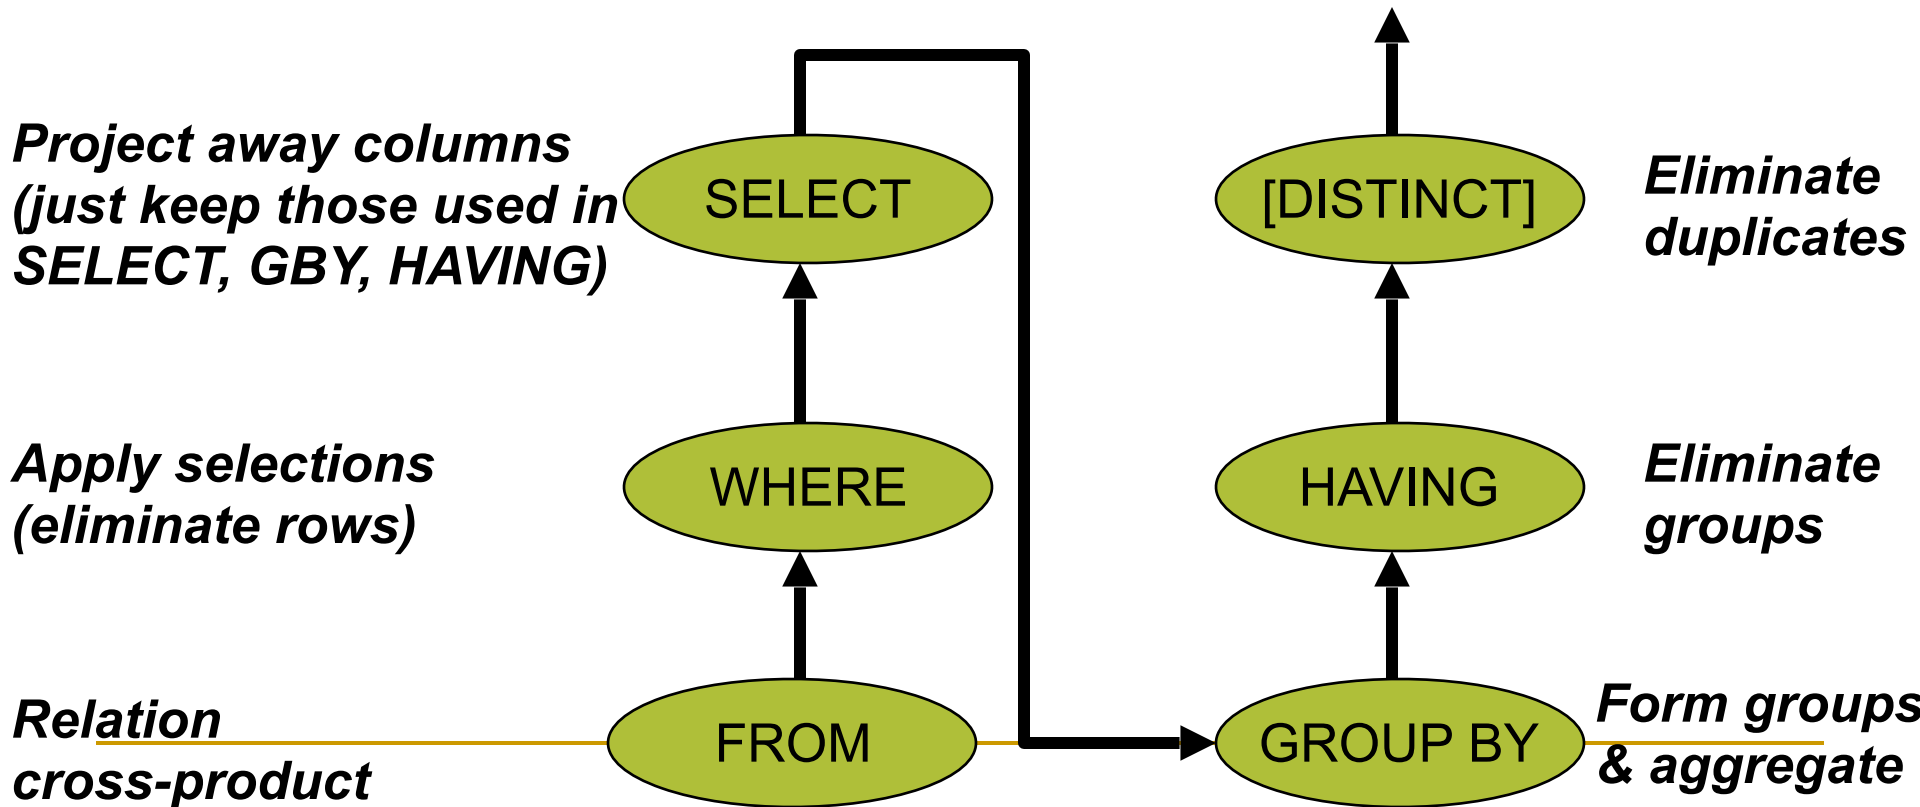
# More on Set-Comparison Operators

- we've seen: **IN, EXISTS**
- can also have: **NOT IN, NOT EXISTS**
- other forms: *op ANY, op ALL*
- Find sailors whose rating is greater than that of some sailor called Horatio:

```
SELECT *  
FROM   Sailors S  
WHERE  S.rating > ANY  
      (SELECT S2.rating  
       FROM   Sailors S2  
       WHERE  S2.sname='Horatio')
```

# Conceptual SQL Evaluation

SELECT	[DISTINCT] <i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>
GROUP BY	<i>grouping-list</i>
HAVING	<i>group-qualification</i>



# Sorting the Results of a Query

- **ORDER BY** *column* [ ASC | DESC ] [, ...]

```
SELECT S.rating, S.sname, S.age
      FROM Sailors S, Boats B, Reserves R
      WHERE S.sid=R.sid
            AND R.bid=B.bid AND B.color='red'
      ORDER BY S.rating, S.sname;
```

- Can order by any column in SELECT list, including expressions or aggs:

```
SELECT S.sid, COUNT (*) AS redrescnt
      FROM Sailors S, Boats B, Reserves R
      WHERE S.sid=R.sid
            AND R.bid=B.bid AND B.color='red'
      GROUP BY S.sid
      ORDER BY redrescnt DESC;
```

# Null Values

- Field values are sometimes *unknown* (e.g., a rating has not been assigned) or *inapplicable* (e.g., no spouse's name).
  - ❑ SQL provides a special value *null* for such situations.
- The presence of *null* complicates many issues. E.g.:
  - ❑ Special operators needed to check if value is/is not *null*.
  - ❑ Is *rating*>8 true or false when *rating* is equal to *null*?  
What about *AND*, *OR* and *NOT* connectives?
  - ❑ We need a 3-valued logic (true, false and *unknown*).

# Joins

```
SELECT (column_list)
FROM table_name
  [INNER | {LEFT | RIGHT | FULL } OUTER] JOIN table_name
    ON qualification_list
WHERE ...
```

Explicit join semantics needed unless it is an INNER join

(INNER is default)

---

---

# Inner Join

Only rows that match the qualification are returned.

```
SELECT s.sid, s.name, r.bid  
FROM Sailors s INNER JOIN Reserves r  
ON s.sid = r.sid
```

Returns only those sailors who have reserved boats.

---



```
SELECT s.sid, s.name, r.bid
FROM Sailors s INNER JOIN Reserves r
ON s.sid = r.sid
```

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103

---

# Left Outer Join

Returns all matched rows, plus all unmatched rows from the table on the left of the join clause

(use nulls in fields of non-matching tuples)

```
SELECT s.sid, s.name, r.bid  
FROM Sailors s LEFT OUTER JOIN Reserves r  
ON s.sid = r.sid
```

---

```
SELECT s.sid, s.name, r.bid
FROM Sailors s LEFT OUTER JOIN Reserves r
ON s.sid = r.sid
```

<u>sid</u>	sname	rating	age
22	Dustin	7	45.0
31	Lubber	8	55.5
95	Bob	3	63.5

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

s.sid	s.name	r.bid
22	Dustin	101
95	Bob	103
31	Lubber	

---

# Right Outer Join

Right Outer Join returns all matched rows, plus all unmatched rows from the table on the right of the join clause

```
SELECT r.sid, b.bid, b.name  
FROM Reserves r RIGHT OUTER JOIN Boats  
    b  
ON r.bid = b.bid
```

---

```
SELECT r.sid, b.bid, b.name
FROM Reserves r RIGHT OUTER JOIN Boats b
ON r.bid = b.bid
```

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

---

# Full Outer Join

Full Outer Join returns all (matched or unmatched) rows from the tables on both sides of the join clause

```
SELECT r.sid, b.bid, b.name  
FROM Reserves r FULL OUTER JOIN Boats b  
ON r.bid = b.bid
```

---

```
SELECT r.sid, b.bid, b.name
FROM Reserves r FULL OUTER JOIN Boats b
ON r.bid = b.bid
```

<u>sid</u>	<u>bid</u>	<u>day</u>
22	101	10/10/96
95	103	11/12/96

<u>bid</u>	bname	color
101	Interlake	blue
102	Interlake	red
103	Clipper	green
104	Marine	red

r.sid	b.bid	b.name
22	101	Interlake
	102	Interlake
95	103	Clipper
	104	Marine

Note: in this case it is the same as the ROJ!

~~bid is a foreign key in reserves, so all reservations must have a corresponding tuple in boats.~~

# Views: Defining External DB Schemas

```
CREATE VIEW view_name  
AS select_statement
```

Makes development simpler

Often used for security

Not “materialized”

```
CREATE VIEW Reds  
AS SELECT B.bid, COUNT (*) AS scount  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```



## Views Instead of Relations in Queries

```
CREATE VIEW Reds  
AS SELECT B.bid, COUNT (*) AS scout  
FROM Boats B, Reserves R  
WHERE R.bid=B.bid AND B.color='red'  
GROUP BY B.bid
```

bid	scout
102	1

Reds

```
SELECT bname, scout  
FROM Reds R, Boats B  
WHERE R.bid=B.bid  
AND scout < 10
```

# Discretionary Access Control

GRANT *privileges* ON *object* TO  
*users* [WITH GRANT OPTION]

- Object can be a **Table** or a **View**
- Privileges can be:
  - Select
  - Insert
  - Delete
  - References (cols) – allow to create a foreign key that references the specified column(s)
  - All
- Can later be REVOKEd
- Users can be single users or groups
- See Chapter 17 for more details.

---

# Two more important topics

- Constraints
- SQL embedded in other languages

# Integrity Constraints (Review)

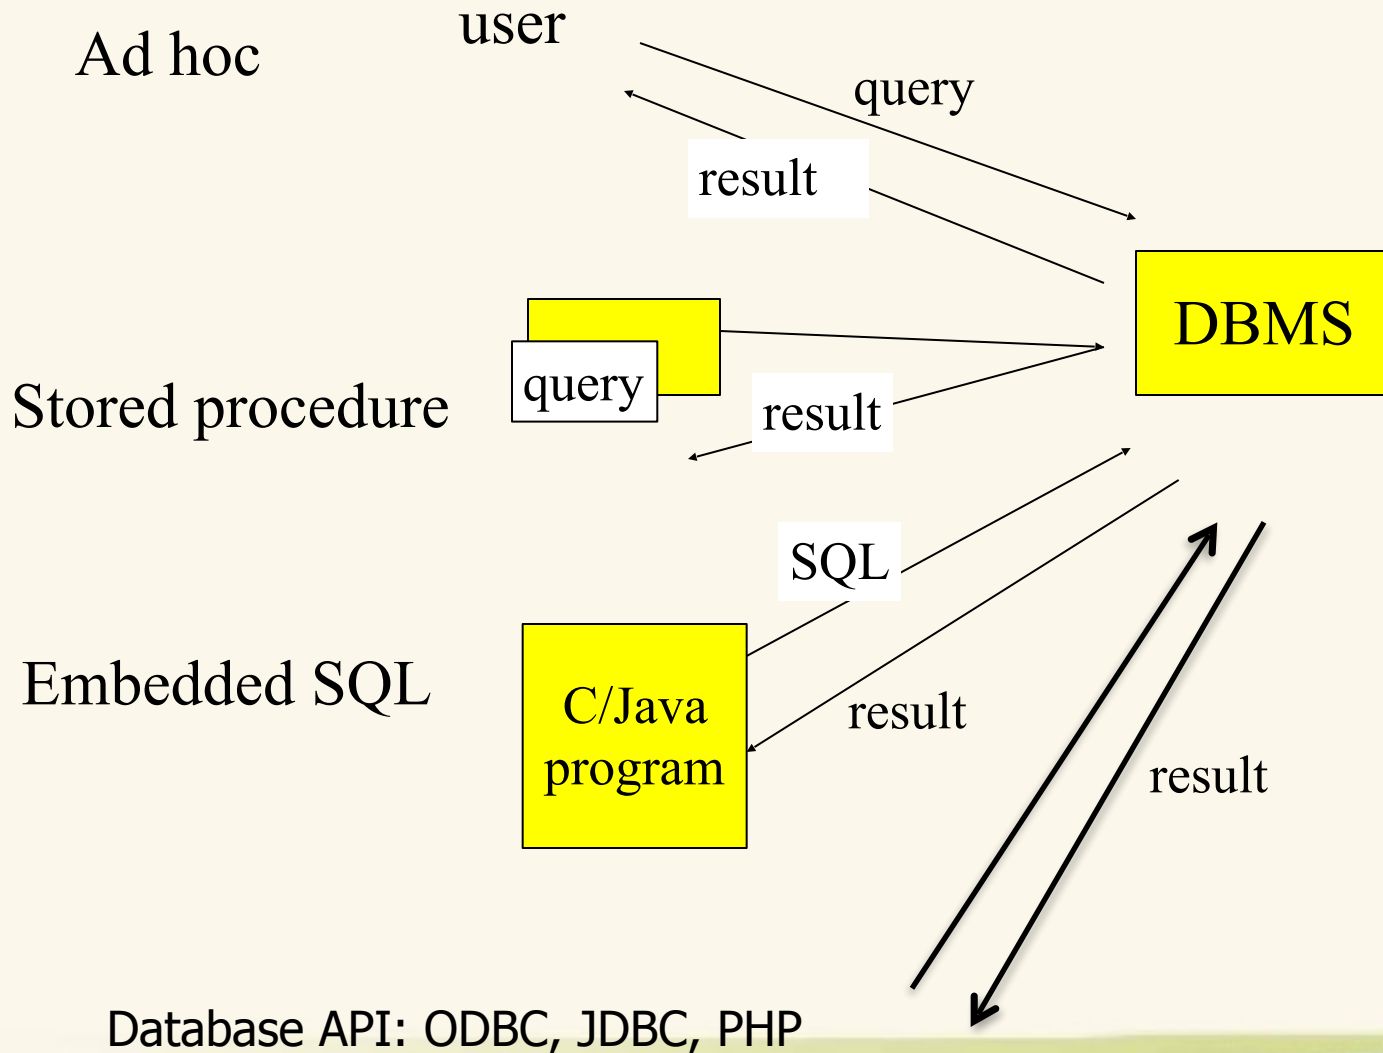
- An IC describes conditions that every *legal instance* of a relation must satisfy.
  - ❑ Inserts/deletes/updates that violate IC's are disallowed.
  - ❑ Can ensure application semantics (e.g., *sid* is a key), or prevent inconsistencies (e.g., *sname* has to be a string, *age* must be  $< 200$ )
- Types of IC's: Domain constraints, primary key constraints, foreign key constraints, general constraints.

# General Constraints

- Useful when more general ICs than keys are involved.
- Can use queries to express constraint.
- Checked on insert or update.
- Constraints can be named.

```
CREATE TABLE Sailors
( sid INTEGER,
  sname CHAR(10),
  rating INTEGER,
  age REAL,
  PRIMARY KEY (sid),
  CHECK ( rating >= 1
        AND rating <= 10 ))
```

```
CREATE TABLE Reserves
( sname CHAR(10),
  bid INTEGER,
  day DATE,
  PRIMARY KEY (bid,day),
  CONSTRAINT noInterlakeRes
  CHECK ('Interlake' <>
        ( SELECT B.bname
          FROM Boats B
          WHERE B.bid=bid)))
```



# Writing Applications with SQL

- SQL is not a general purpose programming language.
  - + Tailored for data retrieval and manipulation
  - + Relatively easy to optimize and parallelize
  - Can't write entire apps in SQL alone

## Options:

Make the query language “Turing complete”

Avoids the “impedance mismatch”

but, loses advantages of relational language simplicity

Allow SQL to be embedded in regular programming languages.

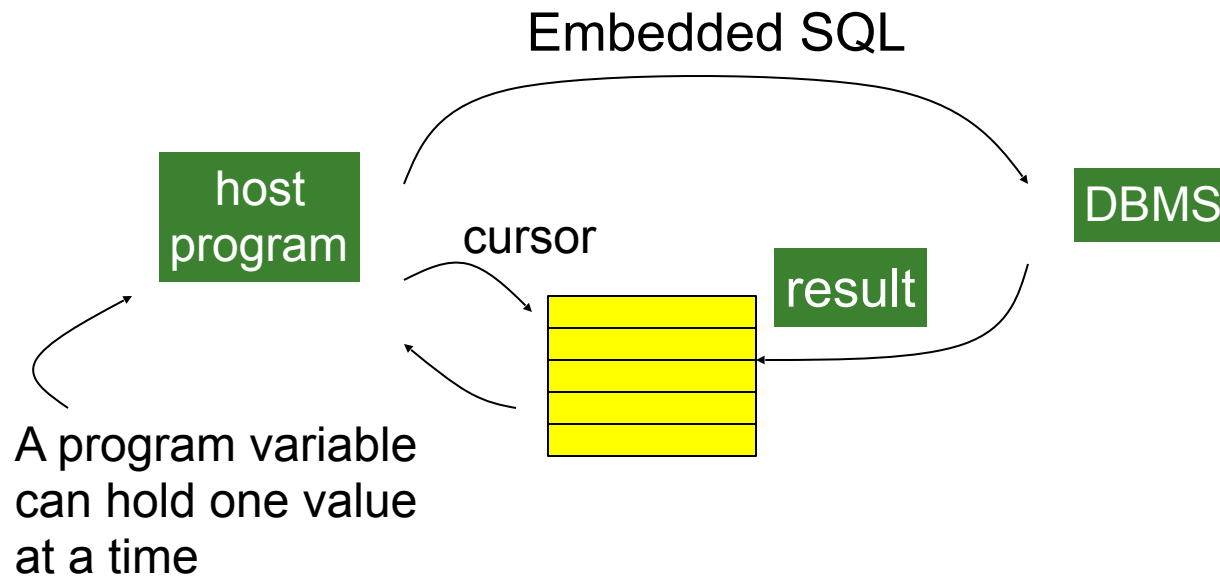
Q: What needs to be solved to make the latter approach work?

# Embedded SQL

- DBMS vendors traditionally provided “host language bindings”
  - E.g. for C or COBOL
  - Allow SQL statements to be called from within a program
  - Typically you preprocess your programs
  - Preprocessor generates calls to a proprietary DB connectivity library
- General pattern
  - One call to *connect* to the right database (login, etc.)
  - SQL statements can refer to *host variables* from the language
- Typically vendor-specific
  - We won't look at any in detail, we'll look at standard stuff
- Problem
  - SQL relations are (multi-)sets, no *a priori* bound on the number of records. No such data structure in C.
  - SQL supports a mechanism called a *cursor* to handle this.



# Why is cursor needed?



Cursor bridges the gap between value-oriented host program and set-oriented DBMS

# Example Embedded SQL

From within a **host** language, find the names and account numbers of customers with more than the variable *amount* dollars in some account.

- Specify the query in SQL and declare a *cursor* for it

**EXEC SQL**

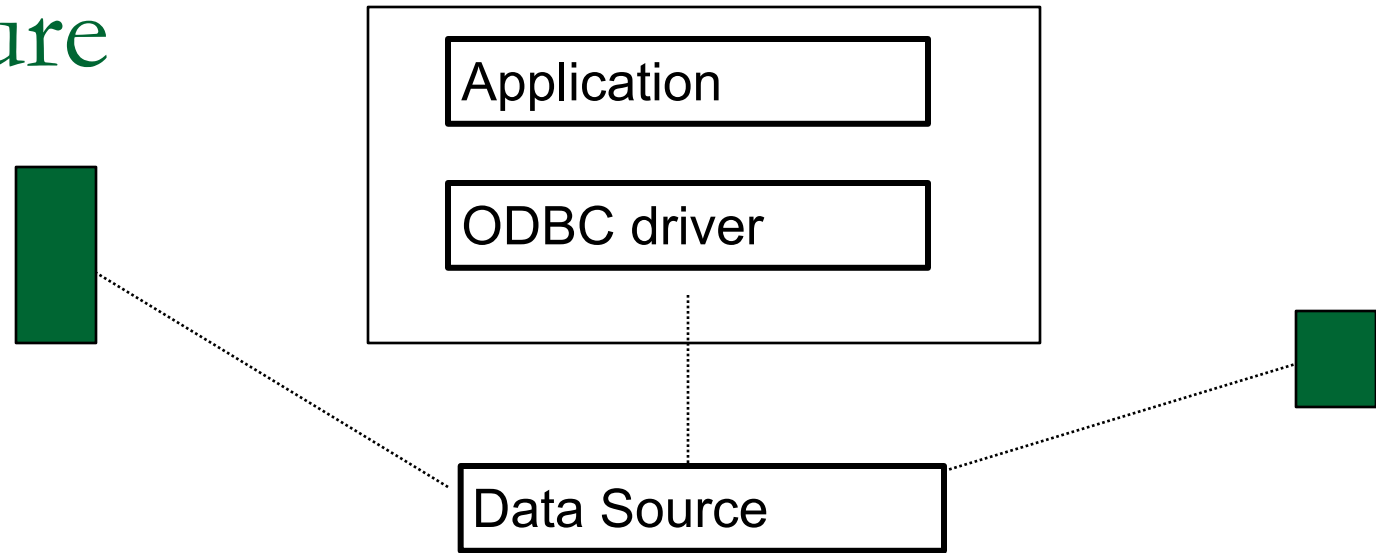
```
declare c cursor for  
select customer-name, account-number  
from depositor, account  
where depositor.account-number = account.account-number and  
account.balance > :amount
```

**END-EXEC**

# Database APIs: Alternative to embedding

- Rather than modify compiler, add a library with database calls (API)
  - ❑ special objects/methods
  - ❑ passes SQL strings from language, presents **result sets** in a language-friendly way
  - ❑ *ODBC* a C/C++ standard started on Windows
  - ❑ *JDBC* a Java equivalent
  - ❑ Most scripting languages have similar things
    - E.g. For Perl there is DBI, “oraPerl”, other packages
- Mostly DBMS-neutral
  - ❑ at least try to hide distinctions across different DBMSs

# Architecture



- A lookup service maps “data source names” (“DSNs”) to drivers
  - Typically handled by OS
- Based on the DSN used, a “driver” is linked into the app at runtime
- The driver traps calls, translates them into DBMS-specific code
- Database can be across a network
- ODBC is standard, so the same program can be used (in principle) to access multiple database systems
- Data source may not even be an SQL database!

# ODBC/JDBC

- Various vendors provide drivers
  - ❑ MS bundles a bunch into Windows
  - ❑ Vendors like DataDirect and OpenLink sell drivers for multiple OSes
- Drivers for various data sources
  - ❑ Relational DBMSs (Oracle, DB2, SQL Server, etc.)
  - ❑ “Desktop” DBMSs (Access, Dbase, Paradox, FoxPro, etc.)
  - ❑ Spreadsheets (MS Excel, Lotus 1-2-3, etc.)
  - ❑ Delimited text files (.CSV, .TXT, etc.)
- You can use JDBC/ODBC *clients* over many data sources
  - ❑ E.g. MS Query comes with many versions of MS Office (msqry32.exe)
- Can write your own Java or C++ programs against xDBC

# JDBC

- Part of Java, easy to use
- Java comes with a JDBC-to-ODBC bridge
  - So JDBC code can talk to any ODBC data source
  - E.g. look in your Windows Control Panel or MacOS Utilities folder for JDBC/ODBC drivers!
- JDBC tutorial online
  - <http://developer.java.sun.com/developer/Books/JDBCTutorial/>

# hello.php

- <http://pages.stern.nyu.edu/~mjohnson/dbms/php/hello.php>

```
<html>
<head><title>Hello from PHP</title>
</head>
<body>
Here comes the PHP part:<BR><BR>
<?php print "Hello, World!<br>\n"; ?>
<br>That's it!
</body></html>
```

- Q: What the difference between `<br>` and `\n`?

# PHP vars

- Names always start with \$
  - <http://pages.stern.nyu.edu/~mjohnson/dbms/php/math.php>

```
<?
$num1 = 58;
$num2 = 67;
print "First number " . $num1 . "<br>";
print "Second number " . $num2 . "<br>";
$total = $num1 + $num2;
print "The sum is " . $total . "<br>";
?>
```



# Combining PHP and HTML

- <http://pages.stern.nyu.edu/~mjohnson/dbms/php/combine.php>

```
<?php
  for($z=0;$z<=5;$z++) {
?>
    Iteration number <? = $z ?><br>
  <?
  }
?>
```

# PHP & MySQL

1. Open a connection and open our DB:

```
$db = mysql_connect("localhost", user, pass);  
mysql_select_db("test", $db);
```

2. Run query:

```
$result = mysql_query($query,$db);
```

# PHP & MySQL

## 3. Extract next row of data from the results:

```
$myrow = mysql_fetch_row($result)
```

- ❑ What this means: myrow is an array that can then be accessed
- ❑ Other options, see code

## ■ In general, to scroll through results, do:

```
while ($myrow = mysql_fetch_row($result))  
    # print row's data
```

# API Summary

**APIs are needed to interface DBMSs to programming languages**

- Embedded SQL uses “native drivers” and is usually faster but less standard
- ODBC (used to be Microsoft-specific) for C/C++
- JDBC the standard for Java
- Scripting languages (PHP, Perl, JSP) are becoming the preferred technique for web-based systems

# Summary

- Relational model has **well-defined query semantics**
- SQL provides functionality close to basic relational model  
*(some differences in duplicate handling, null values, set operators, ...)*
- Typically, many ways to write a query
  - **DBMS figures out a fast way to execute a query, regardless of how it is written.**

# Review

- Examples from [sqlzoo.net](http://sqlzoo.net)

```
SELECT L  
FROM   R1 , ... , Rn  
WHERE  C
```

$$\Pi_L(\sigma_C(R_1 \times \dots \times R_n))$$

# Another complex example

- `People(ssn, name, street, city, state, state)`
- Q: Who lives on George's street?
- A: First, generate pairs of (renamed) people:
  - $\rho_{p1}(\text{People}) \times \rho_{p2}(\text{People})$
- Then pick out pairs with George:
  - $\sigma_{p1.\text{name}='George'}(\rho_{p1}(\text{People}) \times \rho_{p2}(\text{People}))$
- And refine to rows with George and someone else:
  - $\sigma_{p1.\text{name}='George' \text{ AND } p1.\text{name} \neq p2.\text{name}}(\rho_{p1}(\text{People}) \times \rho_{p2}(\text{People}))$
- Finally, project out the names:
  - $\Pi_{p2.\text{name}}(\sigma_{p1.\text{name}='George' \text{ AND } p1.\text{name} \neq p2.\text{name}}(\rho_{p1}(\text{People}) \times \rho_{p2}(\text{People})))$

# Live examples

- Q: produce a list of employees and their bosses
  - ❑ What if no boss? Or no subordinate?
- Joins on **emp**, **emp man**:
  - ❑ Comma-based
  - ❑ Inner
  - ❑ Natural
  - ❑ Cross
  - ❑ Outer – left, right, full



# More live examples

- Inner joins require an **ON** clause
  - Like a where clause
  - Arbitrary boolean expression
  - If always true (**1=1**), reduces to cross join
- New compar op: **BETWEEN**
  - **a between 5 and 10**  $\Leftrightarrow$  **a >= 5 and a <= 10**
- Q: produce a list of employees with their salary grades
  - **emp, salgrade**