


# Relational Query Optimization


**R & G Chapters 12/15**





# Review

- **Implementation of single Relational Operations**
- **Choices depend on **indexes, memory, stats**,...**
- **Joins**
  - **Blocked nested loops:**
    - simple, exploits extra memory
  - **Indexed nested loops:**
    - best if 1 rel small and one indexed
  - **Sort/Merge Join**
    - good with small amount of memory, bad with duplicates
  - **Hash Join**
    - fast (enough memory), bad with skewed data




# Query Optimization Overview

- Query can be converted to relational algebra
- Rel. Algebra converted to tree, joins as branches
- **Each operator has implementation choices**
- **Operators can also be applied in different order!**

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

$\pi_{(sname)} \sigma_{(bid=100 \wedge rating > 5)} (Reserves \bowtie Sailors)$





# Query Optimization Overview (cont.)

- ***Plan:*** *Tree of R.A. ops (and some others) with choice of algorithm for each op.*
- **Three main issues:**
  - For a given query, **what plans are considered?**
  - How is the **cost of a plan estimated?**
  - How do we **“search” in the “plan space”?**
- **Ideally:** Want to find best plan.
- **Reality:** Avoid worst plans!



# Cost-based Query Sub-System

Queries

```
Select *  
From Blah B  
Where B.blah = blah
```

Query Parser

Query Optimizer

Plan  
Generator

Plan Cost  
Estimator

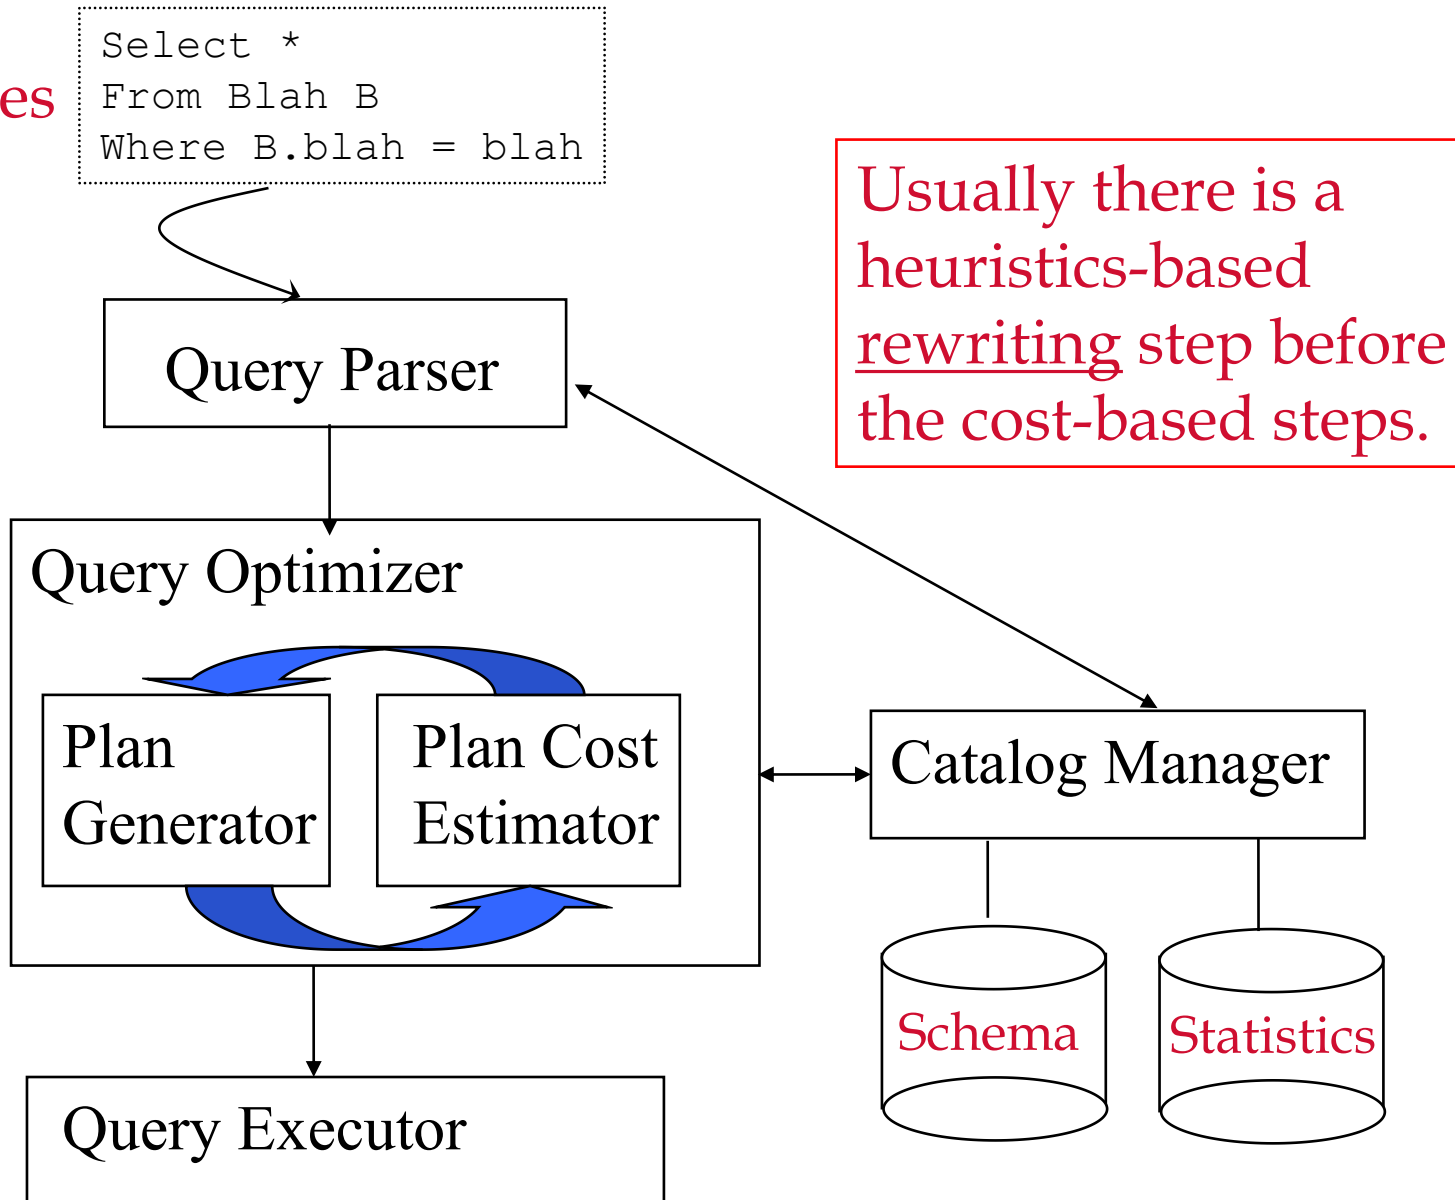
Query Executor

Usually there is a  
heuristics-based  
rewriting step before  
the cost-based steps.

Catalog Manager

Schema

Statistics





# Schema for Examples

Sailors (sid: integer, sname: string, rating: integer, age: real)

Reserves (sid: integer, bid: integer, day: dates, rname: string)

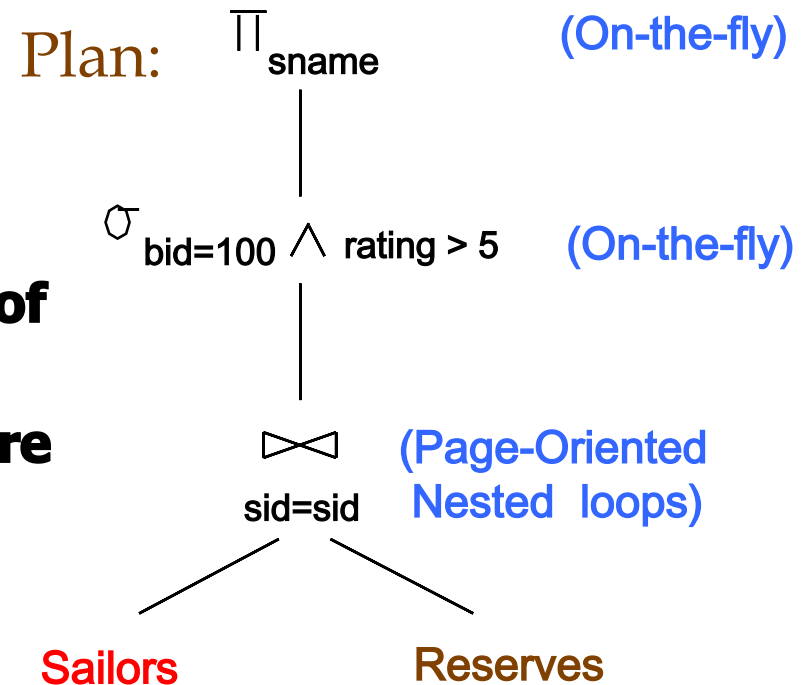
- **As seen in previous lectures...**
- **Reserves:**
  - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
  - Assume there are 100 boats
- **Sailors:**
  - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.
  - Assume there are 10 different ratings
- **Assume we have 5 pages in our buffer pool!**



# Motivating Example

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

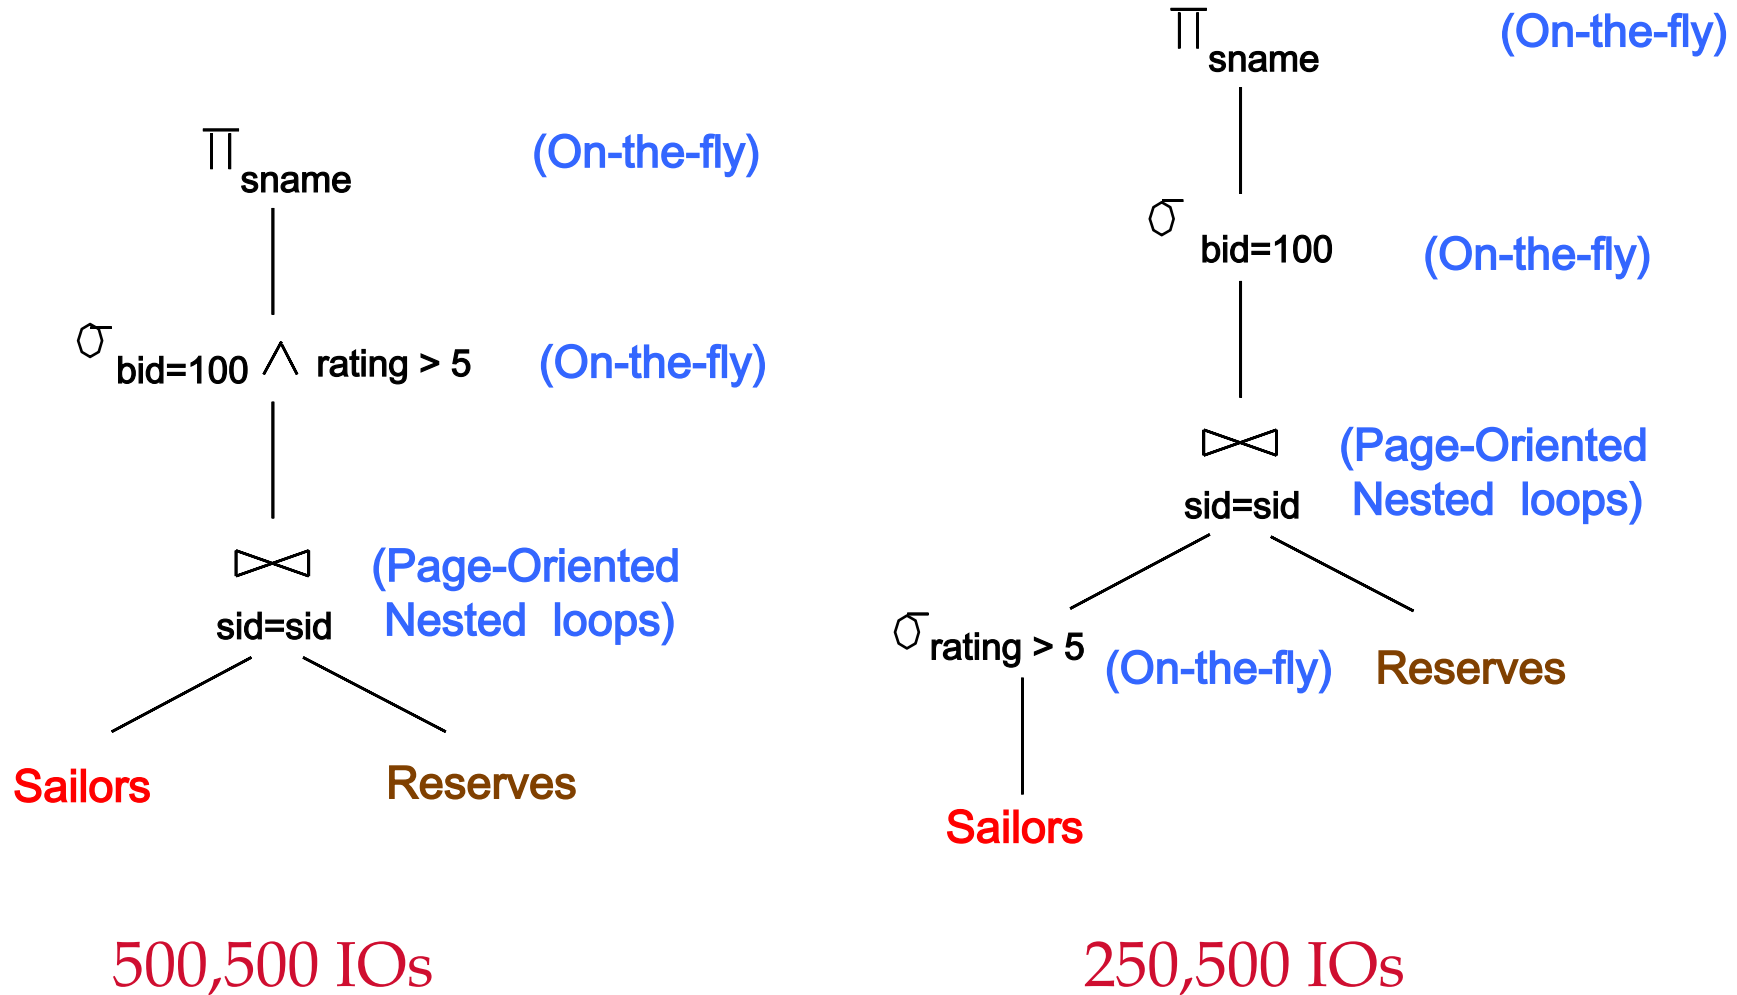
- **Cost: 500+500\*1000 I/Os**
- **By no means the worst plan!**
- **Misses several opportunities:** selections could have been 'pushed' earlier, no use is made of any available indexes, etc.
- **Goal of optimization:** To find more efficient plans that compute the same answer.





# Alternative Plans – Push Selects


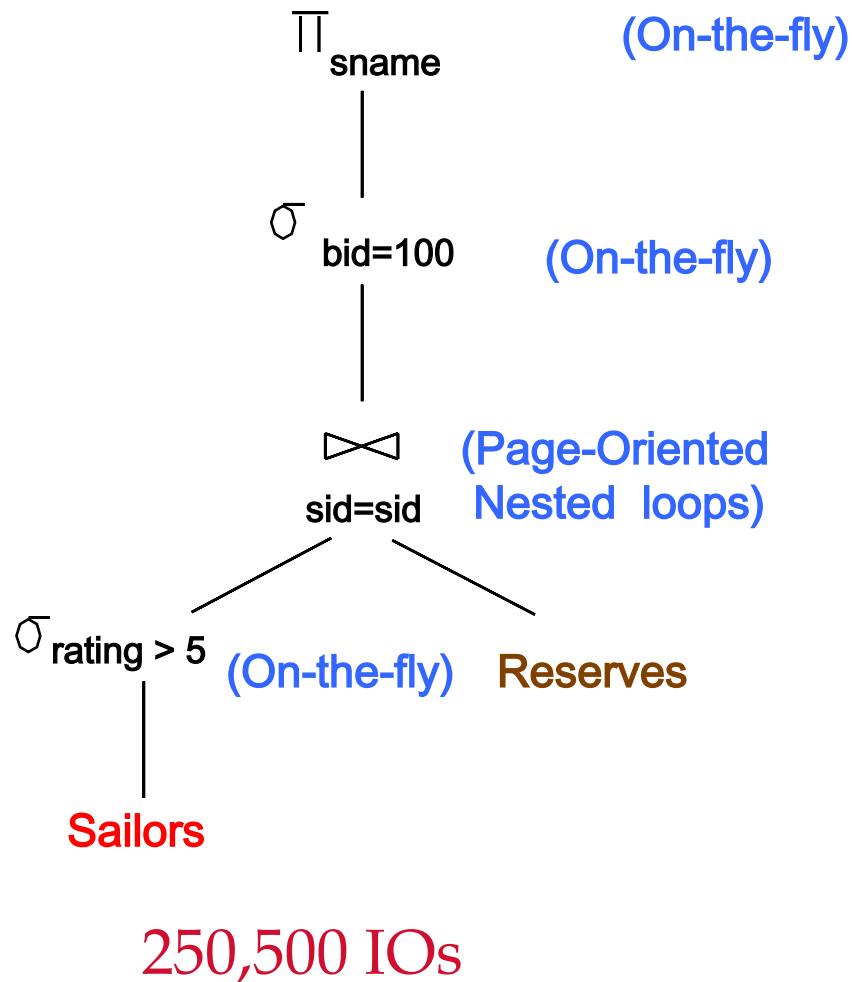
## (No Indexes)





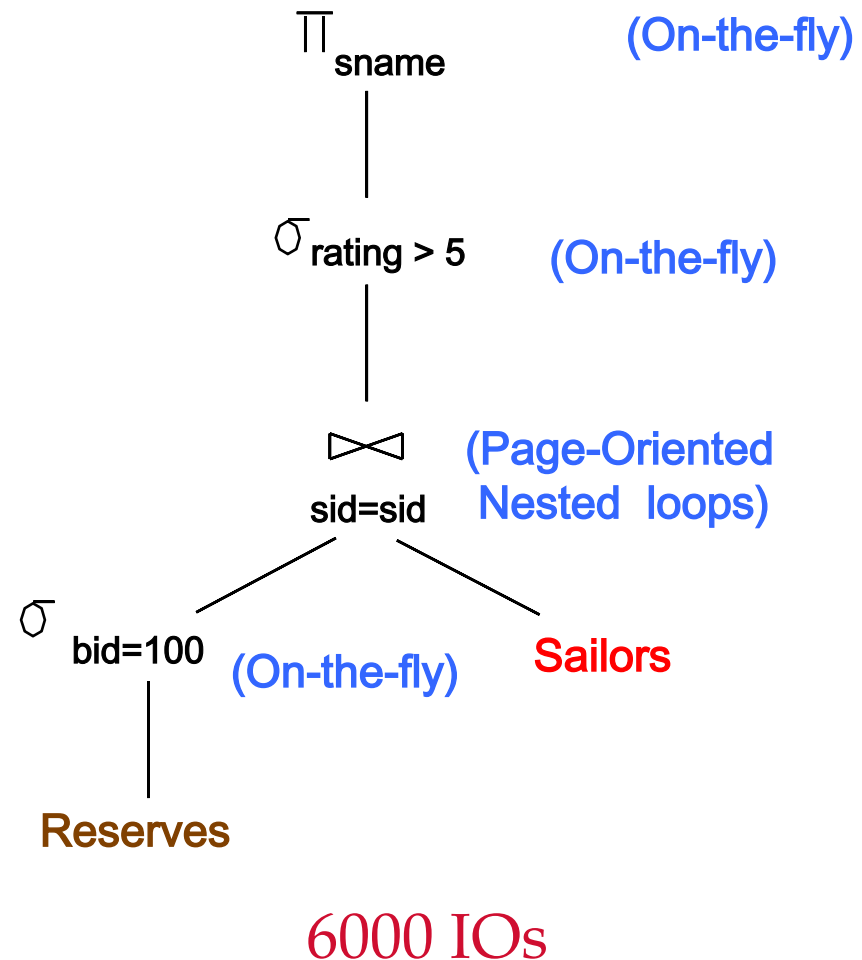
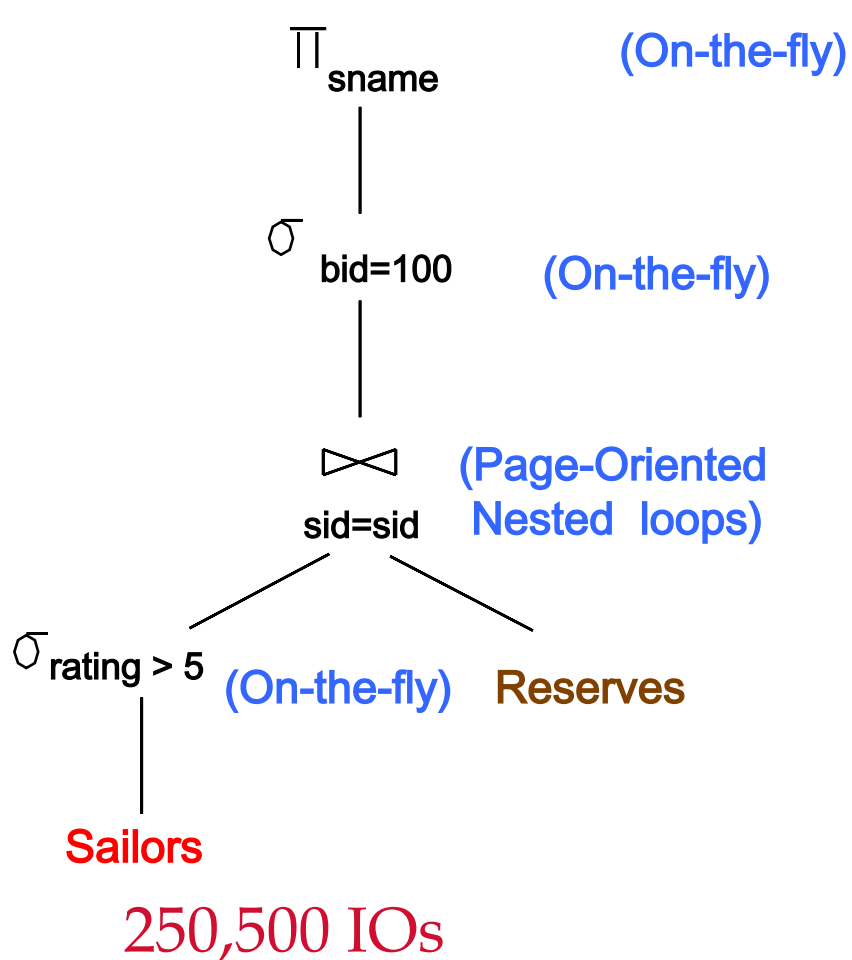


# Alternative Plans – Push Selects (No Indexes)



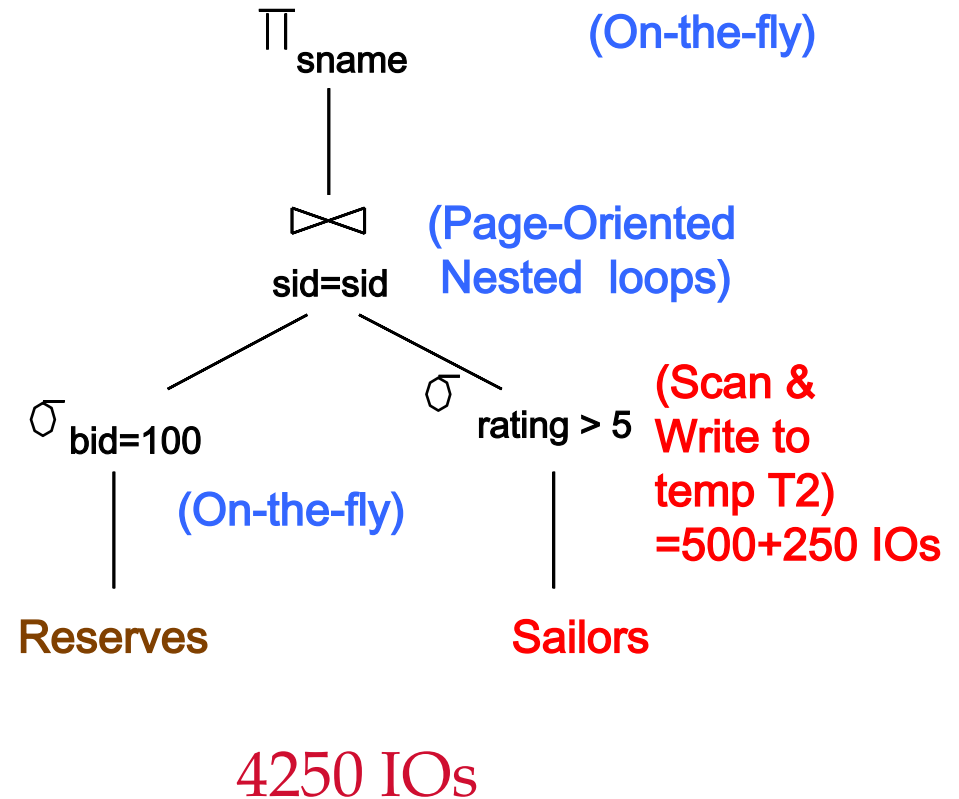
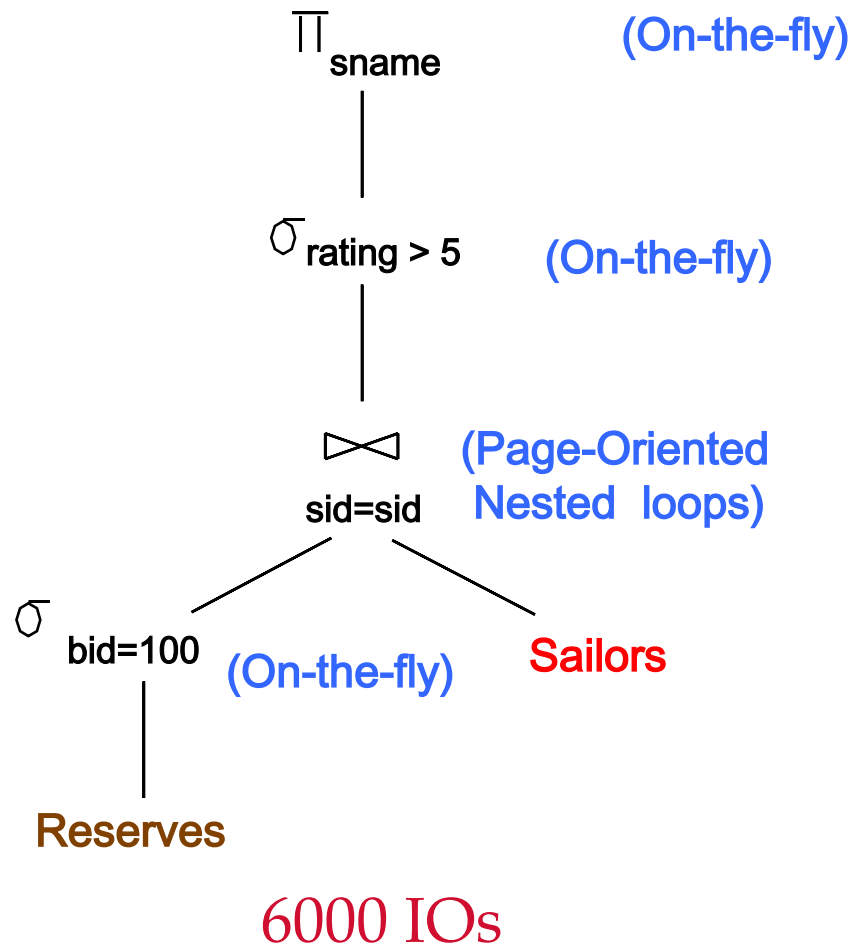



# Alternative Plans – Push Selects (No Indexes)



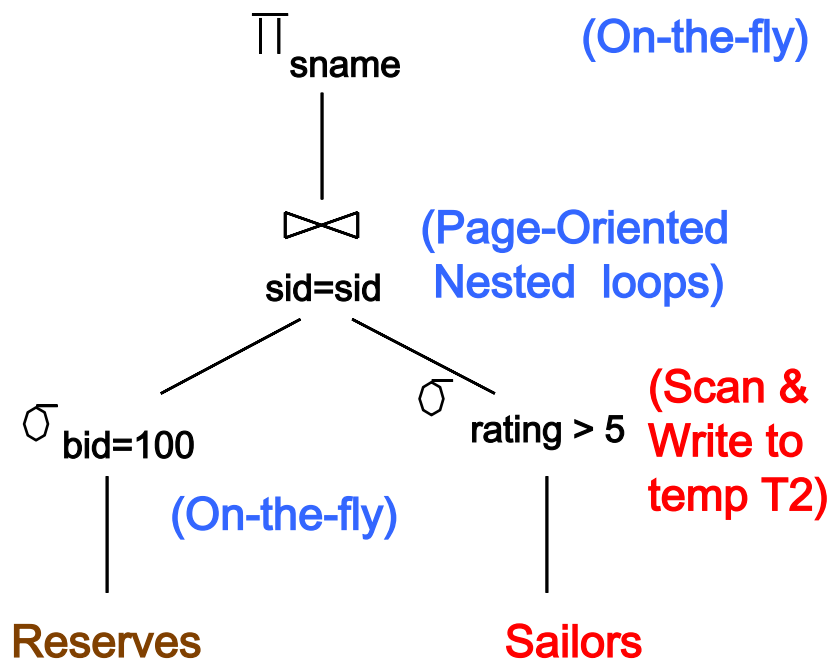


# Alternative Plans – Push Selects (No Indexes)

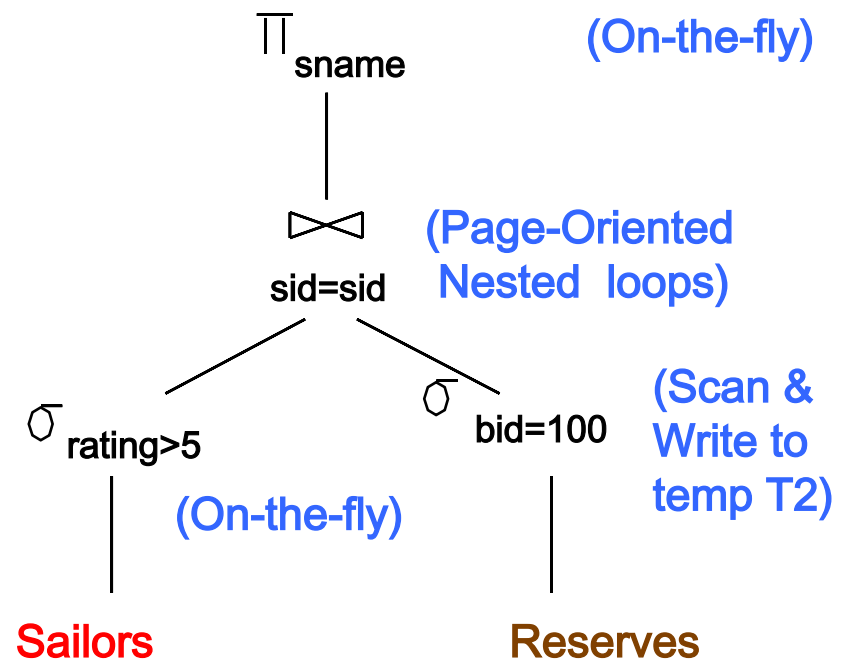




# Alternative Plans – Push Selects (No Indexes)




4250 IOs  
 $= 1000 + 500 + 250 + (10 * 250)$




4010 IOs  
 $= 500 + 1000 + 10 + (250 * 10)$



## More Alternative Plans (No Indexes)



- **Main difference:**  
**Sort Merge Join**
- **With 5 buffers, cost of plan:**
  - Scan Reserves (1000) + write temp T1 (10 pages, if we have 100 boats, uniform distribution) = 1010.
  - Scan Sailors (500) + write temp T2 (250 pages, if have 10 ratings) = 750.
  - Sort T1 ( $2 \times 2 \times 10$ ) + sort T2 ( $2 \times 4 \times 250$ ) + merge (10+250) = 2300
  - **Total: 4060 page I/Os.**
- **If use BNL join, join =  $10 + 4 \times 250$ , total cost = 2770.**
- **Can also 'push' projections, but must be careful!**
  - T1 has only *sid*, T2 only *sid*, *sname*:
  - T1 fits in 3 pgs, cost of BNL under 250 pgs, **total < 2000.**



## More Alt Plans: Indexes

- With clustered index on *bid* of Reserves, we get  $100,000/100 = 1000$  tuples on  $1000/100 = 10$  pages.

- INL with **outer not materialized**.

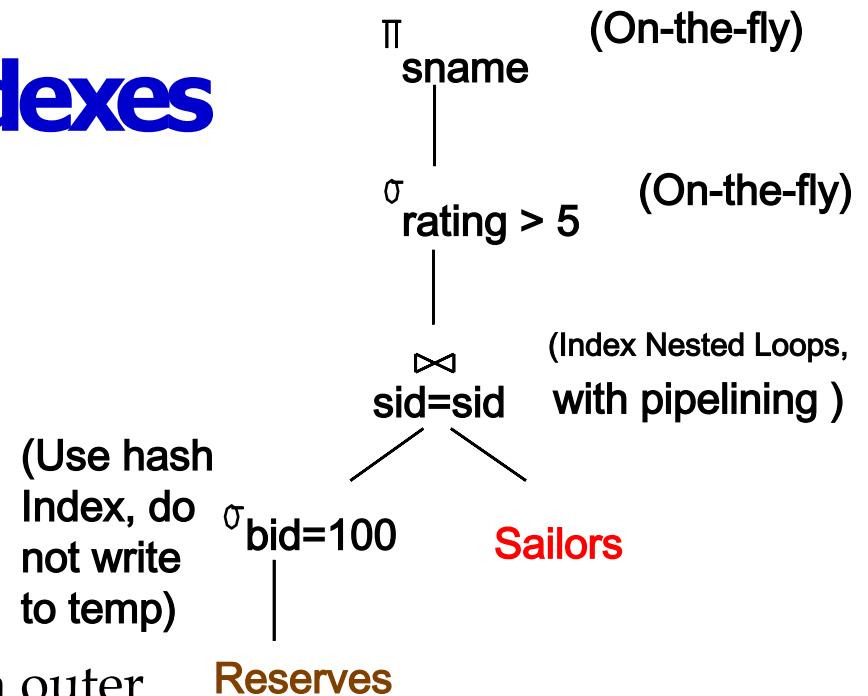
– Projecting out unnecessary fields from outer doesn't help.

- ❖ Join column *sid* is a key for Sailors.

– At most one matching tuple, unclustered index on *sid* OK.

- ❖ Decision not to push *rating*>5 before the join is based on availability of *sid* index on Sailors.

- ❖ **Cost**: Selection of Reserves tuples (10 I/Os); then, for each, must get matching Sailors tuple ( $1000 \times 1.2$ ); total **1210 I/Os**.





# What is needed for optimization?

- **A closed set of operators**
  - Relational ops (table in, table out)
  - Encapsulation (e.g. based on iterators)
- **Plan space**
  - Based on relational equivalences, different implementations
- **Cost Estimation, based on**
  - Cost formulas
  - Size estimation, based on
    - Catalog information on base tables
    - Selectivity (Reduction Factor) estimation
- **A search algorithm**
  - To sift through the plan space based on cost!



# Summary

- **Query optimization is an important task in a relational DBMS.**
- **Must understand optimization in order to understand the performance impact of a given database design (relations, indexes) on a workload (set of queries).**
- **Two parts to optimizing a query:**
  - **Consider a set of alternative plans.**
    - **Must prune search space; typically, *left-deep plans* only.**
  - **Must estimate cost of each plan that is considered.**
    - **Must estimate size of result and cost for each plan node.**
    - ***Key issues:* Statistics, indexes, operator implementations.**





# Query Optimization

- **Query can be dramatically improved by changing access methods, order of operators.**
- **Iterator interface**
- **Cost estimation**
  - **Size estimation and reduction factors**
- **Statistics and Catalogs**
- **Relational Algebra Equivalences**
- **Choosing alternate plans**
- **Multiple relation queries**
- **Will focus on “System R”-style optimizers**



# Highlights of System R Optimizer

- **Impact:**
  - Most widely used currently; works well for  $< 10$  joins.
- **Cost estimation:**
  - Very inexact, but works ok in practice.
  - Statistics, maintained in system catalogs, used to estimate cost of operations and result sizes.
  - Considers combination of CPU and I/O costs.
  - More sophisticated techniques known now.
- **Plan Space:** Too large, must be pruned.
  - Many plans share common, “overpriced” subtrees
    - ignore them all!
  - In some implementations, only the space of *left-deep plans* is considered.
  - Cartesian products avoided in some implementations.



# Query Blocks: Units of Optimization

- Break query into *query blocks*
- Optimized one block at a time
- Uncorrelated nested blocks computed once
- Correlated nested blocks like function calls
  - But sometimes can be “decorrelated”
  - Beyond the scope of introductory course!


```
SELECT S.sname
FROM Sailors S
WHERE S.age IN
    (SELECT MAX (S2.age)
     FROM Sailors S2
     GROUP BY S2.rating)
```

*Outer block*

*Nested block*

❖ For each block, the plans considered are:

- All available access methods, for each relation in FROM clause.
- All *left-deep join trees* (i.e., right branch always a base table, consider all join **orders** and **join methods**.)





# Schema for Examples

Sailors (*sid*: integer, *sname*: string, *rating*: integer, *age*: real)

Reserves (*sid*: integer, *bid*: integer, *day*: dates, *rname*: string)

- **Reserves:**

- Each tuple is 40 bytes long, 100 tuples per page, 1000 pages. 100 distinct bids.

- **Sailors:**

- Each tuple is 50 bytes long, 80 tuples per page, 500 pages. 10 ratings, 40,000 sids.



# Translating SQL to Relational Algebra

```
SELECT S.sid, MIN (R.day)
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = "red"
GROUP BY S.sid
HAVING COUNT (*) >= 2
```

For each sailor with at least two reservations for red boats, find the sailor id and the earliest date on which the sailor has a reservation for a red boat.



# Translating SQL to Relational Algebra

```
SELECT S.sid, MIN (R.day)
FROM Sailors S, Reserves R, Boats B
WHERE S.sid = R.sid AND R.bid = B.bid AND B.color = "red"
GROUP BY S.sid
HAVING COUNT (*) >= 2
```

$\pi$  S.sid, MIN(R.day)  
(HAVING COUNT(\*)>2 (  
GROUP BY S.Sid (  
 $\sigma_{B.color = \text{"red"}}$  (  
Sailors  $\bowtie$  Reserves  $\bowtie$  Boats))))



# Relational Algebra Equivalences

- Allow us to choose different join orders and to 'push' selections and projections ahead of joins.
- Selections:
  - $\sigma_{c1 \wedge \dots \wedge cn}(R) \equiv \sigma_{c1}(\dots(\sigma_{cn}(R))\dots)$  (*cascade*)
  - $\sigma_{c1}(\sigma_{c2}(R)) \equiv \sigma_{c1}(\sigma_{c1}(R))$  (*commute*)
- Projections:
  - $\pi_{a1}(R) \equiv \pi_{a1}(\dots(\pi_{a1, \dots, an}(R))\dots)$  (*cascade*)
- Cartesian Product
  - $R \times (S \times T) \equiv (R \times S) \times T$  (associative)
  - $R \times S \equiv S \times R$  (commutative)
  - *This means we can do joins in any order.*
    - But..beware of cartesian product!



## More Equivalences

- **Eager projection**
  - Can cascade and “push” some projections thru selection
  - Can cascade and “push” some projections below one side of a join
  - Rule of thumb: can project anything not needed “downstream”
- **Selection between attributes of the two arguments of a cross-product converts cross-product to a join.**
- **A selection on just attributes of R commutes with  $R \bowtie S$ . (i.e.,  $\sigma(R \bowtie S) \equiv \sigma(R) \bowtie S$ )**





# Cost Estimation

- For each plan considered, must estimate total cost:
  - Must **estimate cost** of each operation in plan tree.
    - Depends on input cardinalities.
    - We've already discussed how to estimate the cost of operations (sequential scan, index scan, joins, etc.)
  - Must **estimate size of result** for each operation in tree!
    - Use information about the input relations.
    - For selections and joins, assume independence of predicates.
  - In System R, cost is boiled down to a single number consisting of  $\#I/O + \text{factor} * \#CPU \text{ instructions}$
  - Q: Is "cost" the same as estimated "run time"?



# Statistics and Catalogs


- Need information about the relations and indexes involved. *Catalogs* typically contain at least:
  - # tuples (NTuples) and # pages (NPages) per rel'n.
  - # distinct key values (NKeys) for each **index**.
  - low/high key values (Low/High) for each index.
  - Index height (IHeight) for each **tree** index.
  - # index pages (INPages) for each index.
- Catalogs updated periodically.
  - Updating whenever data changes is too expensive; lots of approximation anyway, so slight inconsistency ok.
- More detailed information (e.g., histograms of the values in some field) are sometimes stored.



# Size Estimation and Reduction Factors

```
SELECT attribute list  
FROM relation list  
WHERE term1 AND ... AND termk
```

- **Consider a query block:**
- **Maximum # tuples in result is the product of the cardinalities of relations in the FROM clause.**
- ***Reduction factor (RF)* associated with each *term* reflects the impact of the *term* in reducing result size.**  
***Result cardinality* = Max # tuples \* product of all RF's. (RF = |output| / |input|)**
- ***RF usually called "selectivity"***
  - only R&G seem to call it Reduction Factor
  - beware of confusion between "high selectivity" as defined here and "highly selective" in common English!



# Result Size Estimation

A

1

2

3

4

- *Result cardinality* =  
Max # tuples \* product of all RF's.
- Term *col=value* (given index I on *col*)  
RF =  $1/NKeys(I)$
- Term *col1=col2* (This is handy for joins too...)  
RF =  $1/MAX(NKeys(I1), NKeys(I2))$
- Term *col>value*  
RF =  $(High(I)-value)/(High(I)-Low(I))$

(Implicit assumptions: values are uniformly distributed and terms are independent!)

- Note, if missing indexes, assume 1/10!!!




# Reduction Factors & Histograms

- For better estimation, use a histogram


*equiwidth*

No. of Values	2	3	3	1	8	2	1
Value	0-0.99	1-1.99	2-2.99	3-3.99	4-4.99	5-5.99	6-6.99



*equidepth*


No. of Values	2	3	3	3	3	2	4
Value	0-0.99	1-1.99	2-2.99	3-4.05	4.06-4.67	4.68-4.99	5-6.99





# Think through estimation for joins

- Term *col1=col2*
  - $RF = 1/\text{MAX}(\text{NKeys}(I1), \text{NKeys}(I2))$
- Q: Given a join of R and S, what is the range of possible result sizes (in #of tuples)?
  - If join is on a key for R (and a Foreign Key in S)?
    - A common case, can treat it specially
- General case: join on {A} ({A} is key for neither)
  - estimate each tuple r of R generates  $\text{NTuples}(S)/\text{NKeys}(A,S)$  result tuples, so...
$$\text{NTuples}(\mathbf{R}) * \text{NTuples}(\mathbf{S})/\text{NKeys}(A,\mathbf{S})$$
  - but can also consider it starting with S, yielding:
$$\text{NTuples}(\mathbf{S}) * \text{NTuples}(\mathbf{R})/\text{NKeys}(A,\mathbf{R})$$
  - If these two estimates differ, take the lower one!
    - Q: Why?





# estimation for joins - example

- $\text{NTuples}(\mathbf{R}) * \text{NTuples}(\mathbf{S}) / \text{NKeys}(\mathbf{A}, \mathbf{S})$ 
  - $\text{NTuples}(\mathbf{R}) = 4$ 
    - R有4条记录
  - $\text{NTuples}(\mathbf{S}) = 9$ 
    - S有9条记录
  - $\text{NKeys}(\mathbf{A}, \mathbf{S}) = 3$ 
    - S的A属性共有3种不同的聚值
  - $\text{NTuples}(\mathbf{S}) / \text{NKeys}(\mathbf{A}, \mathbf{S}) = 3$ 
    - 取值1或2或3的记录估计平均有3条
  - **Est # = 4 \* 9 / 3 = 12**
  - **Actual # = 14**


R		S
A	⋈	A
1		1
2		2
3		3
2		2
		3
		3
		2
		2
		2



# Enumeration of Alternative Plans

- **There are two main cases:**
  - **Single-relation plans**
  - **Multiple-relation plans**
- **For queries over a single relation, queries consist of a combination of selects, projects, and aggregate ops:**
  - Each available access path (file scan / index) is considered, and the one with the least estimated cost is chosen.
  - The different operations are essentially carried out together (e.g., if an index is used for a selection, projection is done for each retrieved tuple, and the resulting tuples are *pipelined* into the aggregate computation).





## Cost Estimates for Single-Relation Plans

- **Index I on primary key matches selection:**
  - *Cost is  $Height(I)+1$  for a B+ tree.*
- **Clustered index I matching one or more selects:**
  - *$(NPages(I)+NPages(R)) * \text{product of RF's of matching selects.}$*
- **Non-clustered index I matching one or more selects:**
  - *$(NPages(I)+NTuples(R)) * \text{product of RF's of matching selects.}$*
- **Sequential scan of file:**
  - *$NPages(R).$*

 **Recall:** *Must also charge for duplicate elimination if required*




## Example


```
SELECT S.sid  
FROM Sailors S  
WHERE S.rating=8
```

- If we have an **index on *rating***:
  - Cardinality =  $(1/NKeys(I)) * NTuples(R) = (1/10) * 40000$  tuples
  - **Clustered index**:  $(1/NKeys(I)) * (NPages(I) + NPages(R)) = (1/10) * (50 + 500) = 55$  pages are retrieved. (This is the **cost**)
  - **Unclustered index**:  $(1/NKeys(I)) * (NPages(I) + NTuples(R)) = (1/10) * (50 + 40000) = 401$  pages are retrieved.
- If we have an **index on *sid***:
  - Would have to retrieve all tuples/pages. With a **clustered** index, the **cost** is **50+500**, with **unclustered** index, **50+40000**.
- Doing a **file scan**:
  - We retrieve all file pages (**500**).



- 






# Enumeration of Left-Deep Plans

- Left-deep plans differ only in the order of relations, the access method for each relation, and the join method for each join.
- Enumerated using N passes (if N relations joined):
  - **Pass 1:** Find best 1-relation plan for each relation.
  - **Pass 2:** Find best way to join result of each 1-relation plan (as outer) to another relation. (*All 2-relation plans.*)
  - **Pass N:** Find best way to join result of a (N-1)-relation plan (as outer) to the N'th relation. (*All N-relation plans.*)
- For each subset of relations, retain only:
  - Cheapest plan overall, plus
  - Cheapest plan for each *interesting order* of the tuples.



# The Dynamic Programming Table

Subset of tables in FROM clause	Interesting-order columns	Best plan	Cost
<b>{R, S}</b>	<b>&lt;none&gt;</b>	<b>hashjoin(R, S)</b>	<b>1000</b>
<b>{R, S}</b>	<b>&lt;R.a, S.b&gt;</b>	<b>sortmerge(R,S)</b>	<b>1500</b>



## A Note on “Interesting Orders”

- An intermediate result has an “interesting order” if it is sorted by any of:
  - ORDER BY attributes
  - GROUP BY attributes
  - Join attributes of *yet-to-be-added* (downstream) joins



## Enumeration of Plans (Contd.)

- An N-1 way plan is not combined with an additional relation unless there is a join condition between them, unless all predicates in WHERE have been used up.
  - i.e., **avoid Cartesian products if possible.**
- **ORDER BY, GROUP BY, aggregates** etc. handled as a final step, using either an `interestingly ordered' plan or an additional sort/hash operator.
- In spite of pruning plan space, this approach is **still exponential** in the # of tables.
- Recall that in practice, COST considered is **#IOs + factor \* CPU Inst**



## Example

Sailors:

Hash, B+ on *sid*

Reserves:

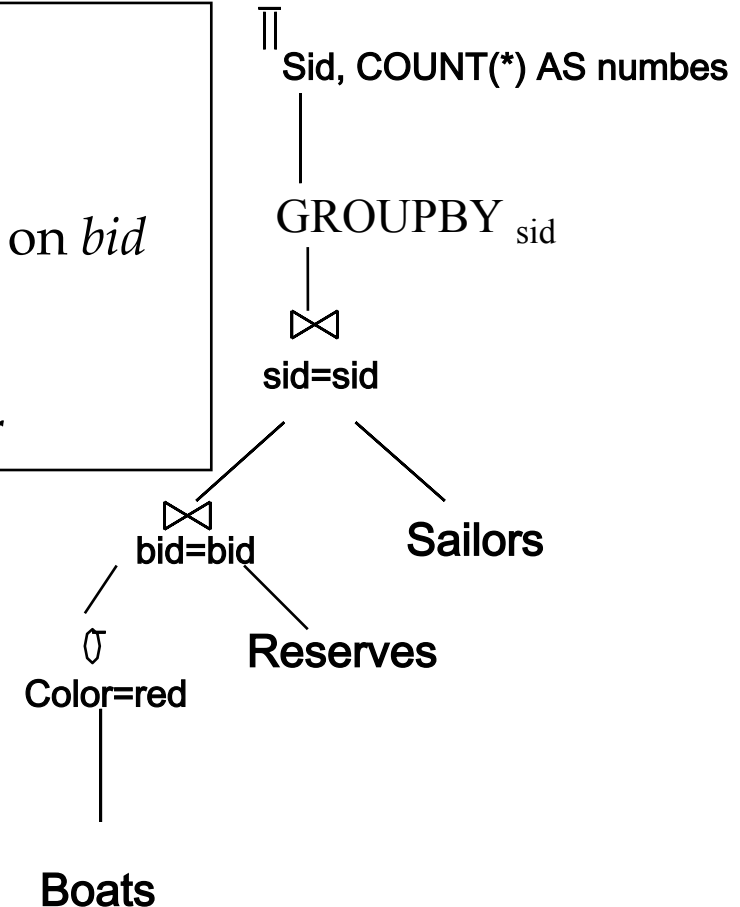
Clustered B+ tree on *bid*

B+ on *sid*

Boats

B+, Hash on *color*

**Select S.sid, COUNT(\*) AS number  
FROM Sailors S, Reserves R, Boats B  
WHERE S.sid = R.sid AND R.bid = B.bid  
AND B.color = "red"  
GROUP BY S.sid**



- **Pass1: Best plan(s) for accessing each relation**
  - Reserves, Sailors: File Scan
  - Q: What about Clustered B+ on Reserves.bid???
  - Boats: B+ tree & Hash on color





## Pass 1

- **Best plan for accessing each relation regarded as the first relation in an execution plan**
  - **Reserves, Sailors: File Scan**
  - **Boats: B+ tree & Hash on color**




## Pass 2

- **For each of the plans in pass 1, generate plans joining another relation as the inner, using all join methods (and matching inner access methods)**
  - File Scan Reserves (outer) with Boats (inner)
  - File Scan Reserves (outer) with Sailors (inner)
  - File Scan Sailors (outer) with Boats (inner)
  - File Scan Sailors (outer) with Reserves (inner)
  - Boats hash on color with Sailors (inner)
  - Boats Btree on color with Sailors (inner)
  - Boats hash on color with Reserves (inner) (sort-merge)
  - Boats Btree on color with Reserves (inner) (BNL)
- **Retain cheapest plan for each pair of relations**




## Pass 3 and beyond

- **For each of the plans retained from Pass 2, taken as the outer, generate plans for the next join**
  - eg Boats hash on color with Reserves (bid) (inner) (sortmerge)  
inner Sailors (B-tree sid) sort-merge
- **Then, add the cost for doing the group by and aggregate:**
  - This is the cost to sort the result by sid, *unless it has already been sorted by a previous operator.*
- **Then, choose the cheapest plan**



## Points to Remember

- **Must understand optimization in order to understand the performance impact of a given database design (relations, indexes) on a workload (set of queries).**
- **Two parts to optimizing a query:**
  - **Consider a set of alternative plans.**
    - Good to prune search space; e.g., left-deep plans only, avoid Cartesian products.
  - **Must estimate cost of each plan that is considered.**
    - Output cardinality and cost for each plan node.
    - *Key issues:* Statistics, indexes, operator implementations.



## Points to Remember

- **Single-relation queries:**
  - All access paths considered, cheapest is chosen.
  - ***Issues:*** Selections that *match* index, whether index key has all needed fields and/or provides tuples in a desired order.



## More Points to Remember

- **Multiple-relation queries:**
  - All single-relation plans are first enumerated.
    - Selections/projections considered as early as possible.
  - Next, for each 1-relation plan, all ways of joining another relation (as inner) are considered.
  - Next, for each 2-relation plan that is `retained`, all ways of joining another relation (as inner) are considered, etc.
  - At each level, for each subset of relations, only best plan for each interesting order of tuples is `retained`.



# Physical DB Design

- **Query optimizer does what it can to use indices, clustering etc.**
- **DataBase Administrator (DBA) is expected to set up physical design well.**
- **Good DBAs understand query optimizers very well.**



## **One Key Decision: Indexes**

- **Which tables**
- **Which field(s) should be the search key?**
- **Multiple indexes?**
- **Clustering?**






# Index Selection

- **One approach:**
  - Consider most important queries in turn.
  - Consider best plan using the current indexes
  - See if better plan is possible with an additional index.
  - If so, create it
- **But consider impact on updates!**
  - Indexes can make queries go faster, updates slower.
  - Require disk space, too.



## Issues to Consider in Index Selection


- **Attributes mentioned in a WHERE clause are candidates for index search keys.**
  - Range conditions are sensitive to clustering
  - Exact match conditions don't require clustering
    - Or do they???? :-)
- **Choose indexes that benefit many queries**
- **NOTE: only one index can be clustered per relation!**
  - So choose it wisely!



## Example 1

```
SELECT E.ename, D.mgr  
FROM Emp E, Dept D  
WHERE E.dno=D.dno AND D.dname='Toy'
```

- **B+ tree index on *D.dname* supports 'Toy' selection.**
  - Given this, index on *D.dno* is not needed.
- **B+ tree on *E.dno* allows us to get matching (inner) Emp tuples for each selected (outer) Dept tuple.**
- **What if WHERE included: `` ... AND E.age=25'' ?**
  - Could retrieve Emp tuples using index on *E.age*, then join with Dept tuples satisfying *dname* selection.
    - Comparable to strategy that used *E.dno* index.
  - So, if *E.age* index is already created, this query provides much less motivation for adding an *E.dno* index.



## Example 2

```
SELECT E.ename, D.mgr  
FROM Emp E, Dept D  
WHERE E.sal BETWEEN 10000 AND 20000  
AND E.hobby='Stamps' AND E.dno=D.dno
```

- **All selections are on Emp so it should be the outer relation in any Index NL join.**
  - Suggests that we build a B+ tree index on *D.dno*.
- **What index should we build on Emp?**
  - B+ tree on *E.sal* could be used, OR an index on *E.hobby* could be used.
  - Only one of these is needed, and which is better depends upon the selectivity of the conditions.
    - As a rule of thumb, equality selections more selective than range selections.
- ***Have to understand optimizers to get this right!***



## Examples of Clustering

- **B+ tree index on *E.age* can be used to get qualifying tuples.**

- How selective is the condition?
- Is the index clustered?

```
SELECT E.dno  
FROM Emp E  
WHERE E.age>40
```


- **Consider the GROUP BY query.**

```
SELECT E.dno, COUNT (*)  
FROM Emp E  
WHERE E.age>10  
GROUP BY E.dno
```

- If many tuples have *E.age* > 10, using *E.age* index and sorting the retrieved tuples may be costly.
- Clustered *E.dno* index may be better!

- **Equality queries and duplicates:**
  - Clustering on *E.hobby* helps!

```
SELECT E.dno  
FROM Emp E  
WHERE E.hobby=Stamps
```



# Summary

- **Optimization is the reason for the lasting power of the relational system**
- **But it is primitive in some ways**
- **New areas: Smarter summary statistics (fancy histograms and “sketches”), auto-tuning statistics, adaptive runtime re-optimization (e.g. *eddies*)**