



Concurrency Control


R&G - Chapter 17



Smile, it is the key that fits the
lock of everybody's heart.


Anthony J. D'Angelo,
The College Blue Book






Review

- **ACID transaction semantics.**
- **Today: focus on Isolation property**
 - Serial schedules safe but slow
 - Try to find schedules *equivalent* to serial ...



Conflicting Operations

- Need a tool to decide if 2 schedules are equivalent
- Use notion of “conflicting” operations
- Definition: Two operations **conflict** if:
 - They are by different transactions,
 - they are on the same object,
 - and at least one of them is a write.



Conflict Serializable Schedules

- **Definition:** Two schedules are **conflict equivalent** iff:
 - They involve the same actions of the same transactions, and
 - every pair of conflicting actions is ordered the same way
- **Definition:** Schedule S is **conflict serializable** if:
 - S is conflict equivalent to some serial schedule.
- **Note, some “serializable” schedules are NOT conflict serializable**
 - A price we pay to achieve efficient enforcement.

| T1 | T2 | T3 |
|------|------|------|
| R(A) | | |
| | W(A) | |
| W(A) | | |
| | | W(A) |



Conflict Serializability – Intuition

- A schedule **S** is conflict serializable if:
 - You are able to transform **S** into a serial schedule by swapping consecutive non-conflicting operations of different transactions.
- *Example:*

| | | |
|---------------------|---------------------|-----------|
| R(A) W(A) | | R(B) W(B) |
| | R(A) W(A) | R(B) W(B) |
| ≡ | | |
| R(A) W(A) R(B) W(B) | | |
| | R(A) W(A) R(B) W(B) | |




Conflict Serializability (Continued)

- Here's another example:


$R(A)$ $W(A)$
 $R(A)$ $W(A)$

- Serializable or not????


NOT!



Dependency Graph




- **Dependency graph:**
 - One node per Xact
 - Edge from T_i to T_j if:
 - An operation O_i of T_i conflicts with an operation O_j of T_j and
 - O_i appears earlier in the schedule than O_j .
- **Theorem:** Schedule is conflict serializable *if and only if* its dependency graph is acyclic.



Example

- **A schedule that is not conflict serializable:**

| | | |
|-----|------------------------|------------|
| T1: | R(A), W(A), | R(B), W(B) |
| T2: | R(A), W(A), R(B), W(B) | |



- **The cycle in the graph reveals the problem. The output of T1 depends on T2, and vice-versa.**



Two-Phase Locking (2PL)

Lock
Compatibility
Matrix


| | S | X |
|---|---|---|
| S | ✓ | — |
| X | — | — |

rules:

- Xact must obtain a **S** (*shared*) lock before reading, and an **X** (*exclusive*) lock before writing.
- **Xact cannot get new locks after releasing any locks.**



Two-Phase Locking (2PL), cont.



2PL guarantees conflict serializability



But, does not prevent Cascading Aborts.






Strict 2PL


- ***Problem:*** Cascading Aborts
- ***Example:*** rollback of T1 requires rollback of T2!


| | | |
|-----|-------------|-------------------|
| T1: | R(A), W(A), | R(B), W(B), Abort |
| T2: | R(A), W(A) | |

- **Strict Two-phase Locking (Strict 2PL) protocol:**
Same as 2PL, except:
Locks released only when transaction completes
i.e., either:
 (a) transaction has committed (commit record on disk),
 or
 (b) transaction has aborted and rollback is complete.




Strict 2PL (continued)





Next ...


- **A few examples**



Non-2PL, A= 1000, B=2000, Output =?




| | |
|-------------------|-------------------|
| Lock_X(A) | |
| Read(A) | Lock_S(A) |
| A: = A-50 | |
| Write(A) | |
| Unlock(A) | |
| | Read(A) |
| | Unlock(A) |
| | Lock_S(B) |
| Lock_X(B) | |
| | Read(B) |
| | Unlock(B) |
| | PRINT(A+B) |
| Read(B) | |
| B := B +50 | |
| Write(B) | |
| Unlock(B) | |



2PL, A= 1000, B=2000, Output =?



| | |
|-------------------|-------------------|
| Lock_X(A) | |
| Read(A) | Lock_S(A) |
| A: = A-50 | |
| Write(A) | |
| Lock_X(B) | |
| Unlock(A) | |
| | Read(A) |
| | Lock_S(B) |
| | |
| Read(B) | |
| B := B +50 | |
| Write(B) | |
| Unlock(B) | Unlock(A) |
| | Read(B) |
| | Unlock(B) |
| | PRINT(A+B) |



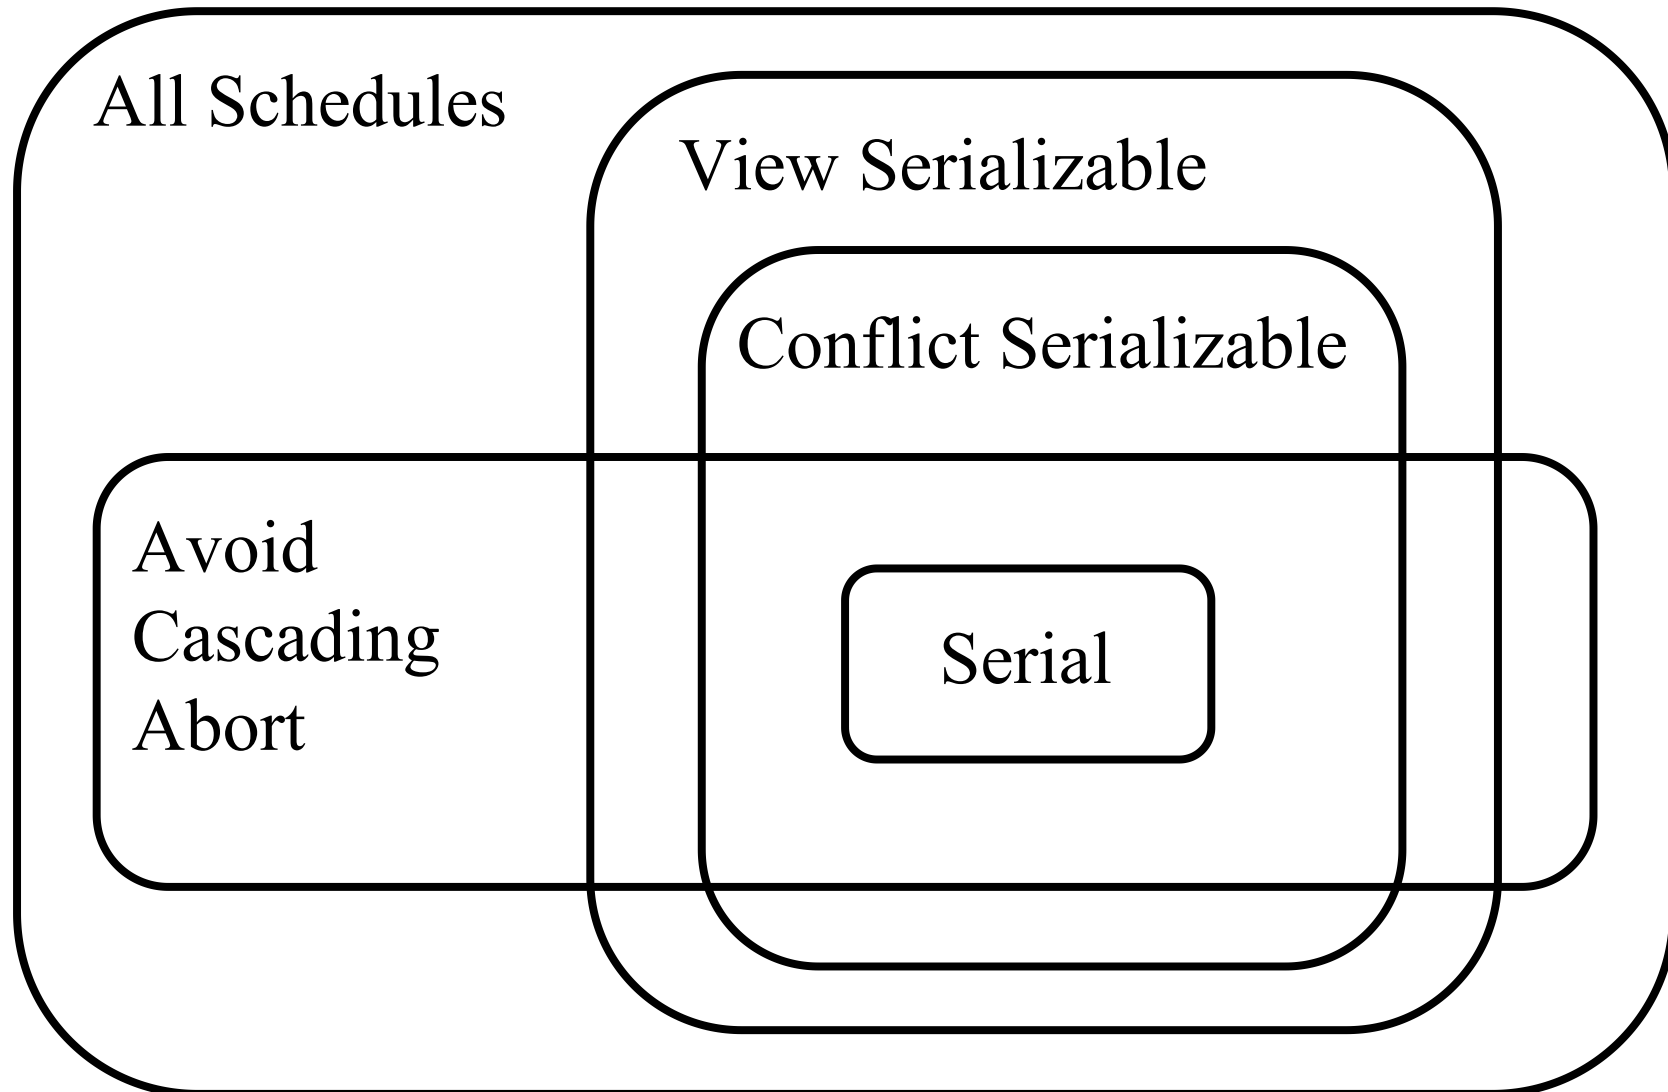
Strict 2PL, $A = 1000$, $B = 2000$, Output = ?



| | |
|--------------------|-------------------|
| Lock_X(A) | |
| Read(A) | Lock_S(A) |
| A := A - 50 | |
| Write(A) | |
| Lock_X(B) | |
| Read(B) | |
| B := B + 50 | |
| Write(B) | |
| Unlock(A) | |
| Unlock(B) | |
| | Read(A) |
| | Lock_S(B) |
| | Read(B) |
| | PRINT(A+B) |
| | Unlock(A) |
| | Unlock(B) |

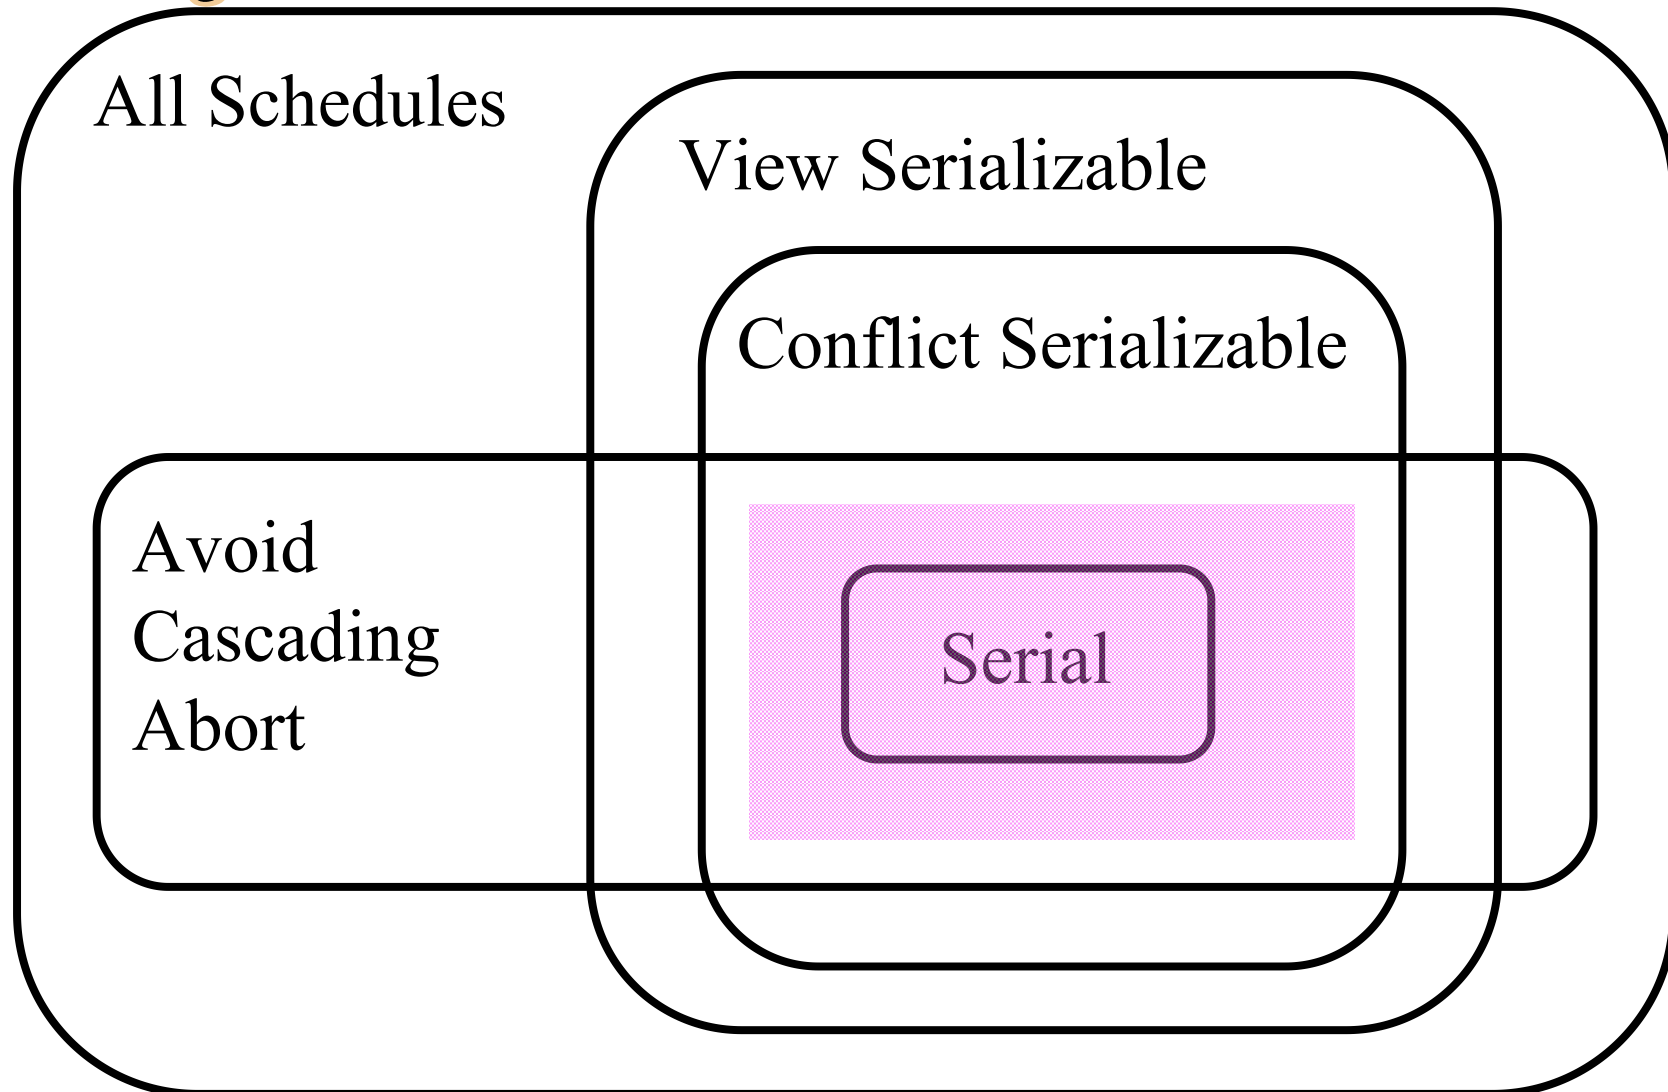


Venn Diagram for Schedules






Which schedules does Strict 2PL allow?





Lock Management

- Lock and unlock requests handled by Lock Manager
- LM keeps an entry for each currently held lock.
- Entry contains:
 - List of xacts currently holding lock
 - Type of lock held (shared or exclusive)
 - Queue of lock requests



Lock Management, cont.

- **When lock request arrives:**
 - Does any other xact hold a conflicting lock?
 - If no, grant the lock.
 - If yes, put requestor into wait queue.
- **Lock upgrade:**
 - xact with shared lock can request to upgrade to exclusive

[illegible]



Deadlocks

- **Deadlock:** Cycle of transactions waiting for locks to be released by each other.
- **Two ways of dealing with deadlocks:**
 - **prevention**
 - **detection**
- **Many systems just punt and use Timeouts**
 - **What are the dangers with this approach?**



Deadlock Detection


- Create and maintain a “waits-for” graph
- Periodically check for cycles in graph



Deadlock Detection (Continued)

Example:

T1: S(A), S(D), S(B)
T2: X(B) X(C)
T3: S(D), S(C), X(A)
T4: X(B)





Deadlock Prevention

- Assign priorities based on timestamps.
- Say T_i wants a lock that T_j holds

Two policies are possible:

Wait-Die: If T_i has higher priority, T_i waits for T_j ;
otherwise T_i aborts

Wound-wait: If T_i has higher priority, T_j aborts;
otherwise T_i waits

- Why do these schemes guarantee no deadlocks?
- Important detail: If a transaction re-starts, make sure it gets its original timestamp. -- Why?



Summary

- **Correctness criterion for isolation is “serializability”.**
 - In practice, we use “conflict serializability,” which is somewhat more restrictive but easy to enforce.
- **Two Phase Locking and Strict 2PL: Locks implement the notions of conflict directly.**
 - The lock manager keeps track of the locks issued.
 - **Deadlocks** may arise; can either be prevented or detected.